

[S06] Introduction to P4

Programming Protocol Independent Packet Processors

Redes Software
Múltiples Grados


Curso 2024-2025

Based on slides provided by Pablo Molinero for the Master on SDN/NFV of UC3M

P4

Contents

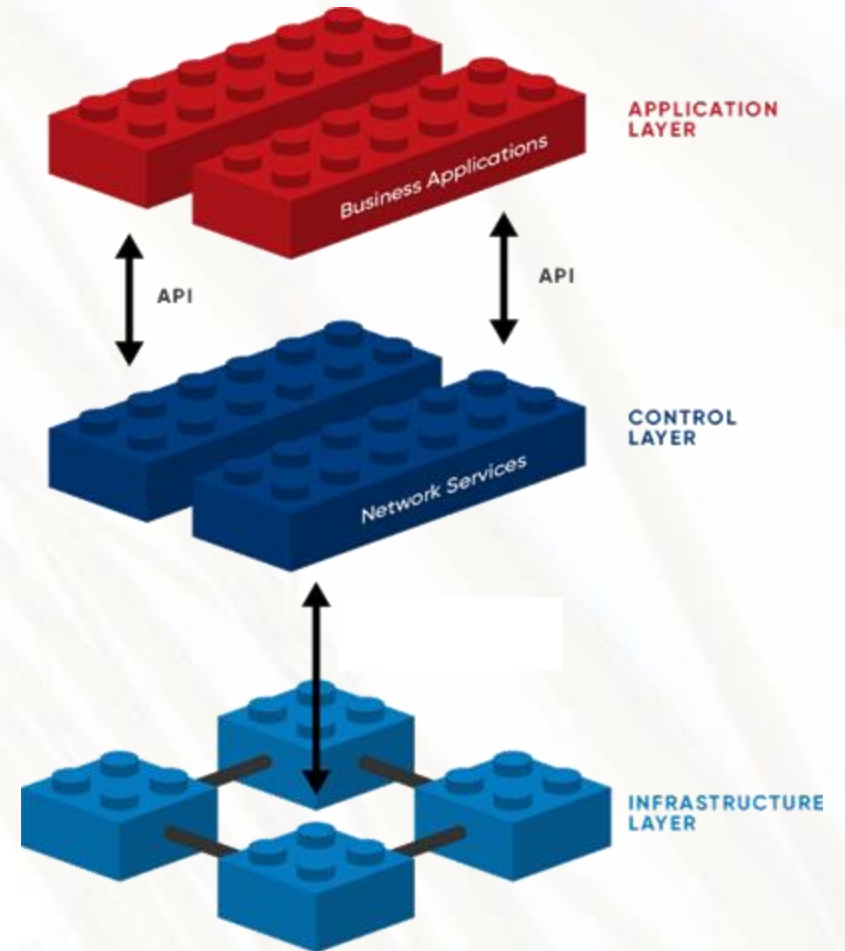
- Why P4?
- What is P4?
- Standards & History
- Use Cases
- P4 Overview
- P4 Language Elements
- Discussion




P4

Contents

- Why P4?
- What is P4?
- Standards & History
- Use Cases
- P4 Overview
- P4 Language Elements
- Discussion



Isn't Open-Flow Enough?



Open-Flow version

Open-flow has *never* been enough: it keeps changing to describe new protocols

Why P4?


- What if I want to create a new protocol
 - Takes years to standardize it
 - High-performance implementations determined by chip manufacturing
 - E.g., VXLAN routing in Broadcom's chips
- The P4 Language came as way to build Advanced switch architectures
 - Originally described in a 2014 SIGCOMM paper titled *"Programming Protocol-Independent Packet Processors"*
 - Programmable packet headers
 - Stateful packet processing

What P4 Brings

- Consequences of P4
 - HW performance with SW programmability
 - Device manufacturer \neq device programmer
 - Many network capabilities exposed to software
- We can innovate in these areas without waiting for the device manufacturing
 - In-band Network Telemetry (INT)
 - Sub-flow load balancer
 - Service chaining
- Good for programmable switches, as well as fixed-function ones

Change of Paradigm

Current: Fixed-Function Switching Chip

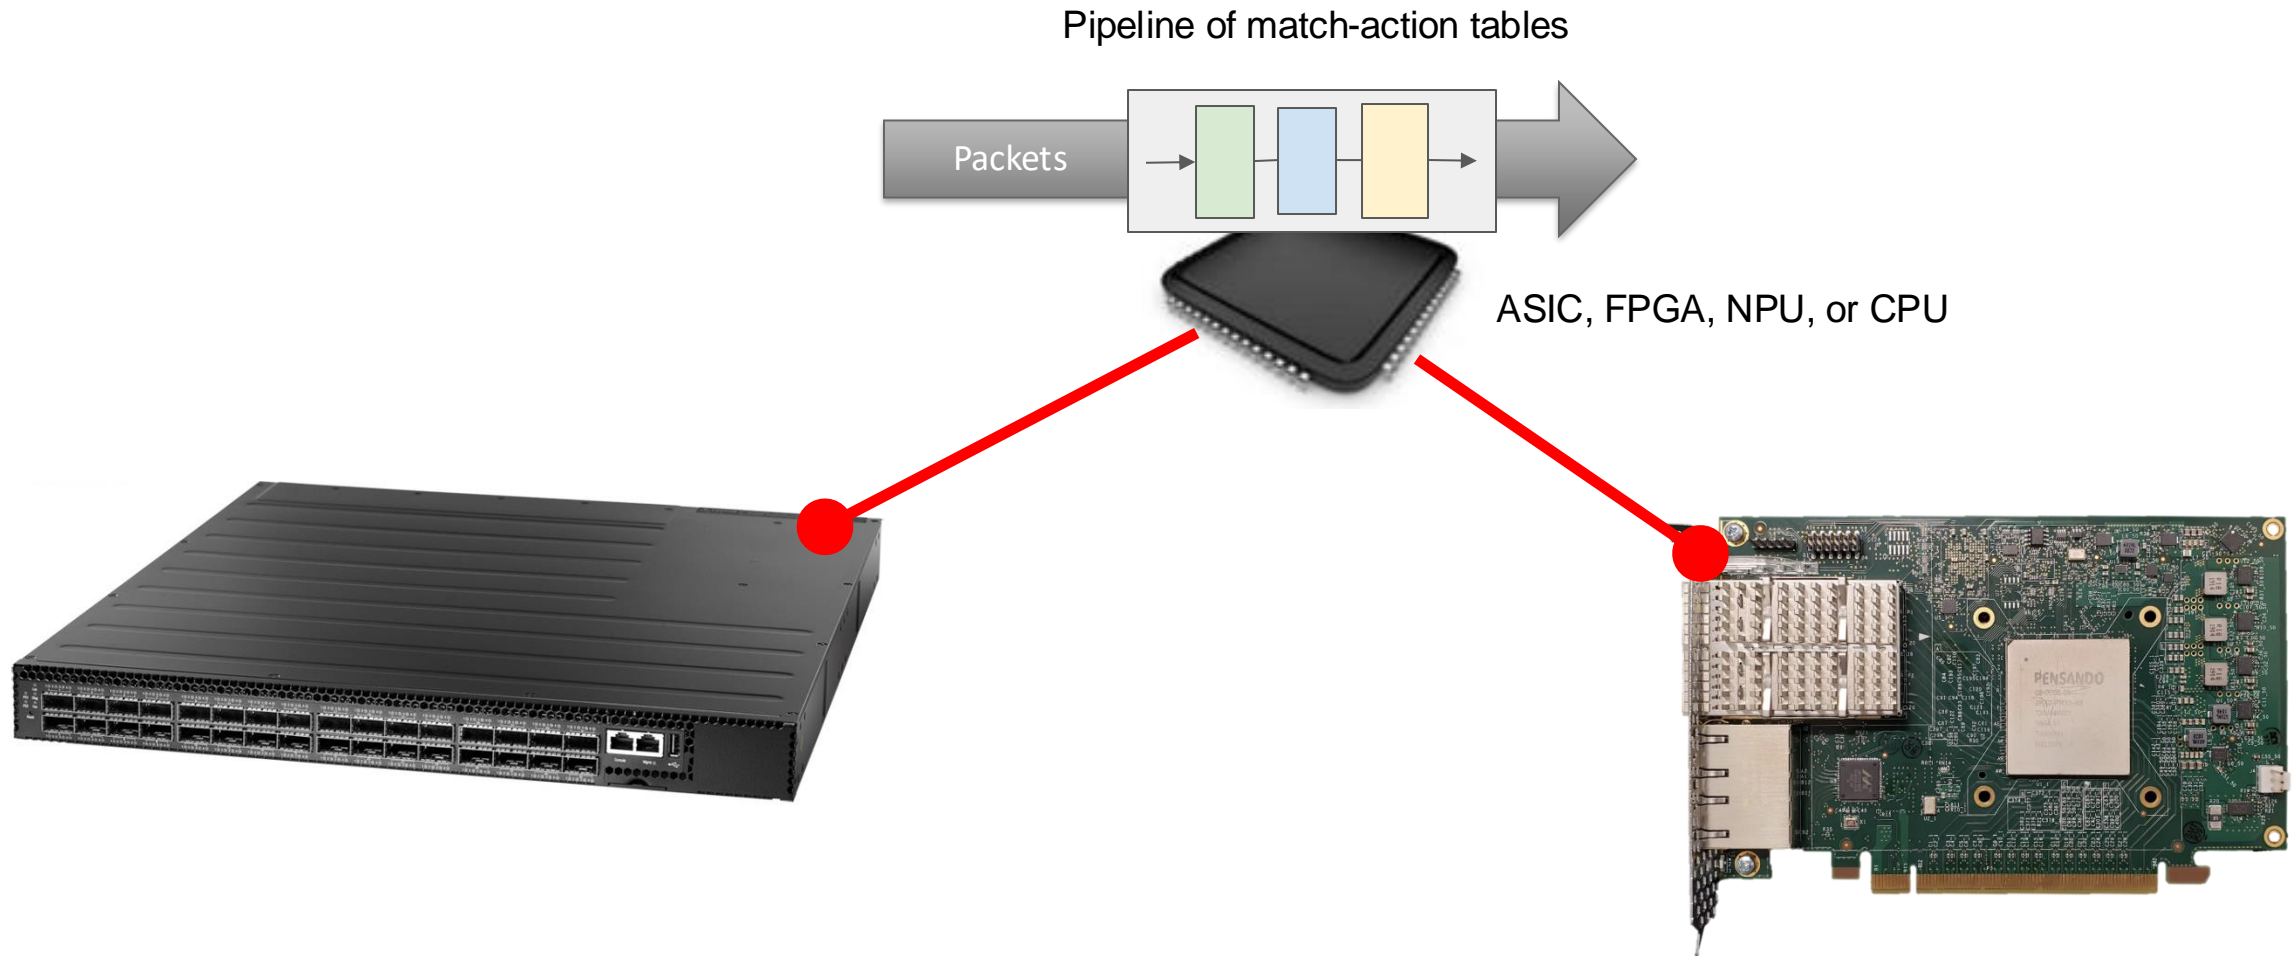


Future: Programmable Switching Chip



Custom
protocol


P4 in a Switch and SmartNIC



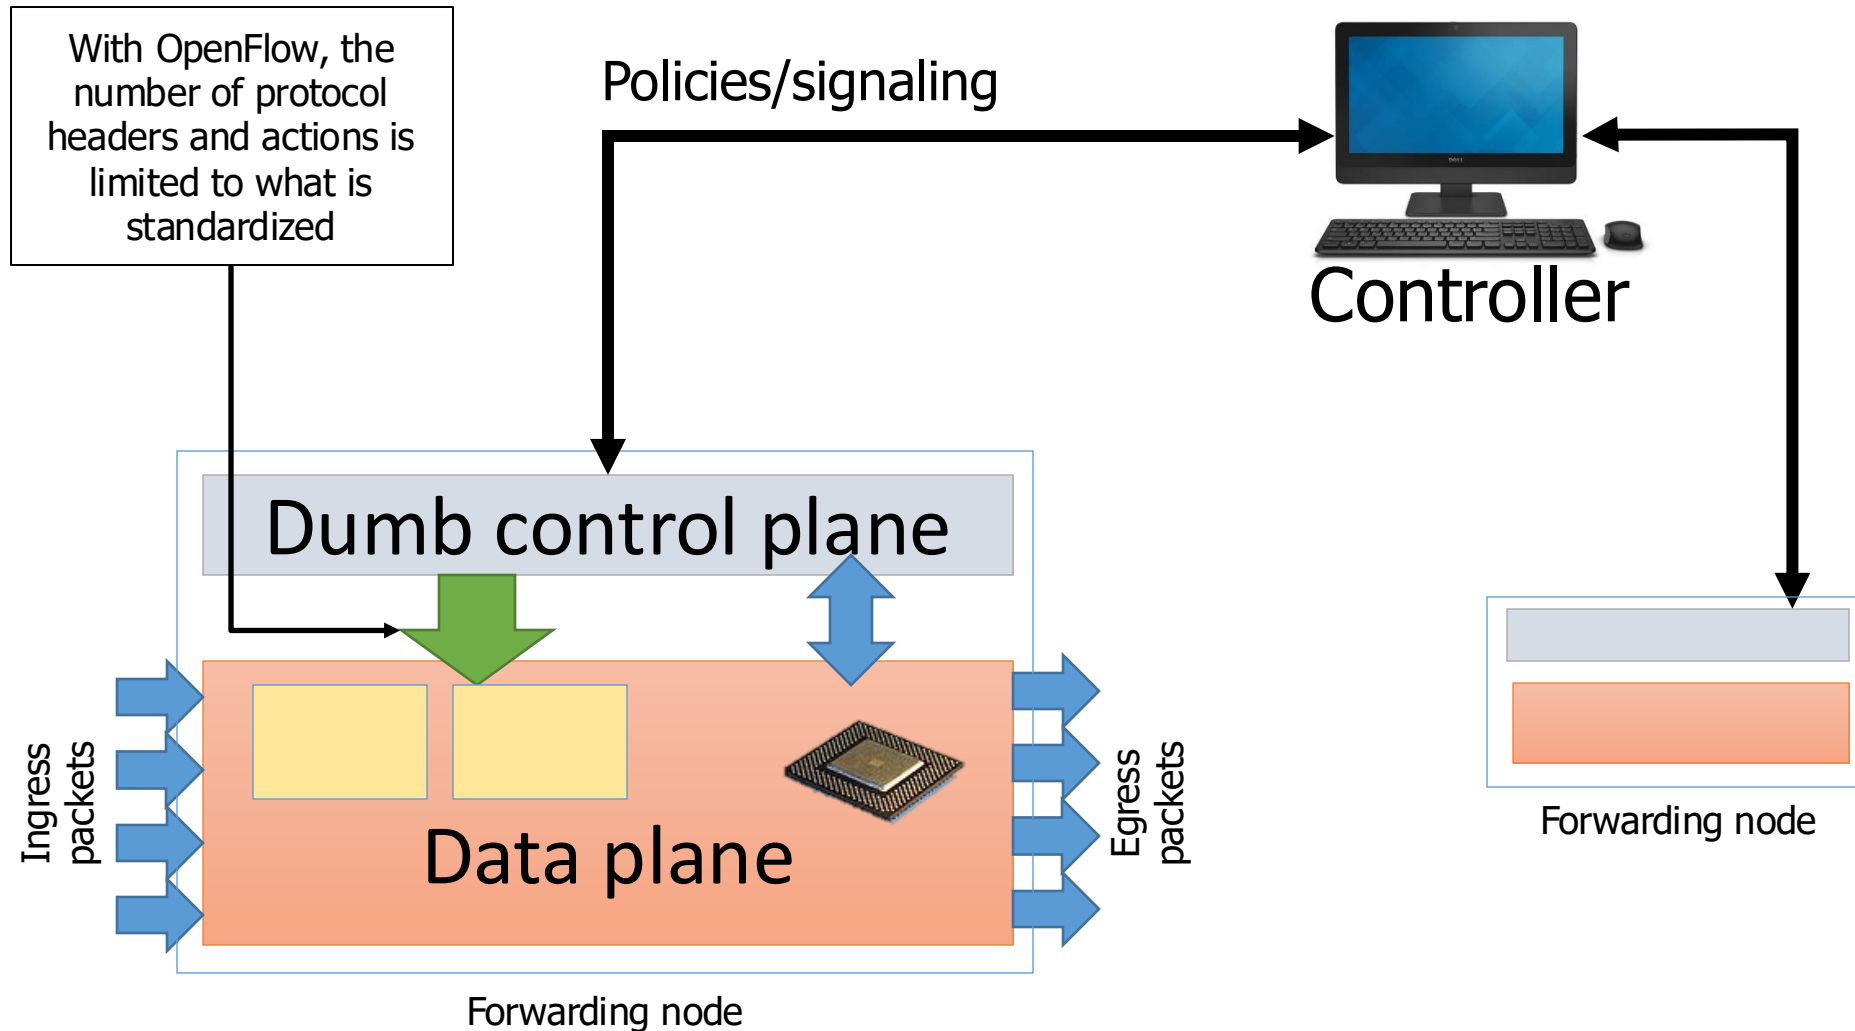
P4

Contents

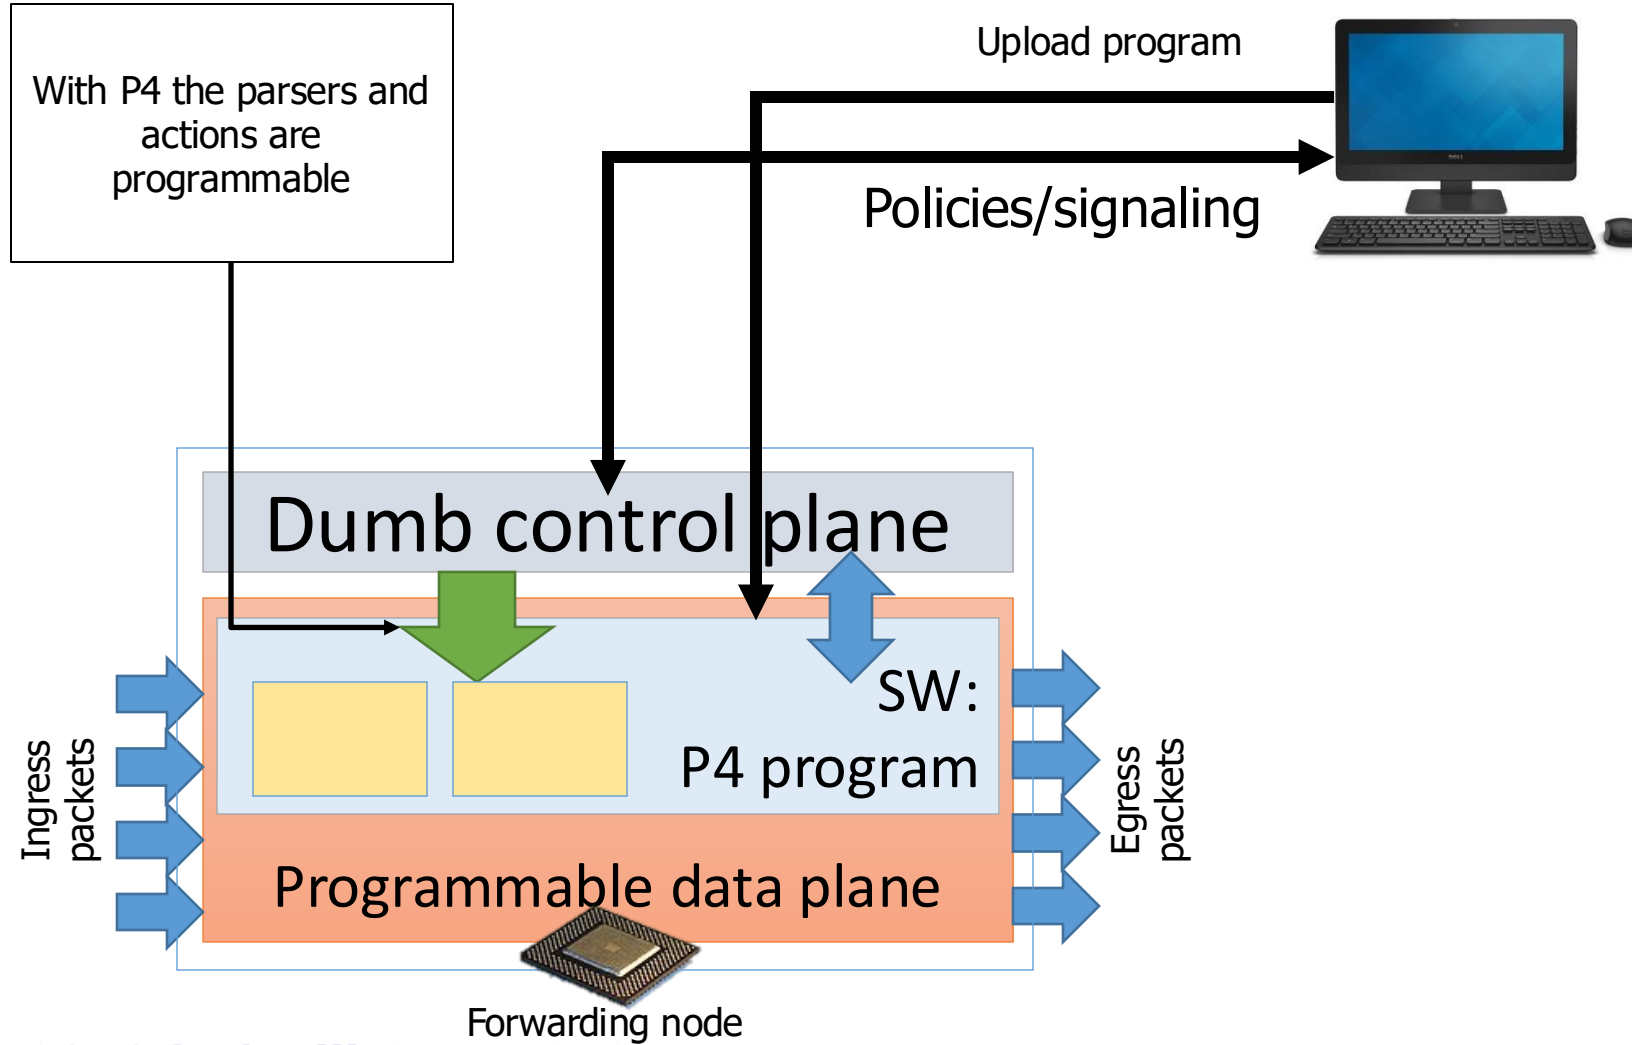
- Why P4?
- What is P4?
- Standards & History
- Use Cases
- P4 Overview
- P4 Language Elements
- Discussion



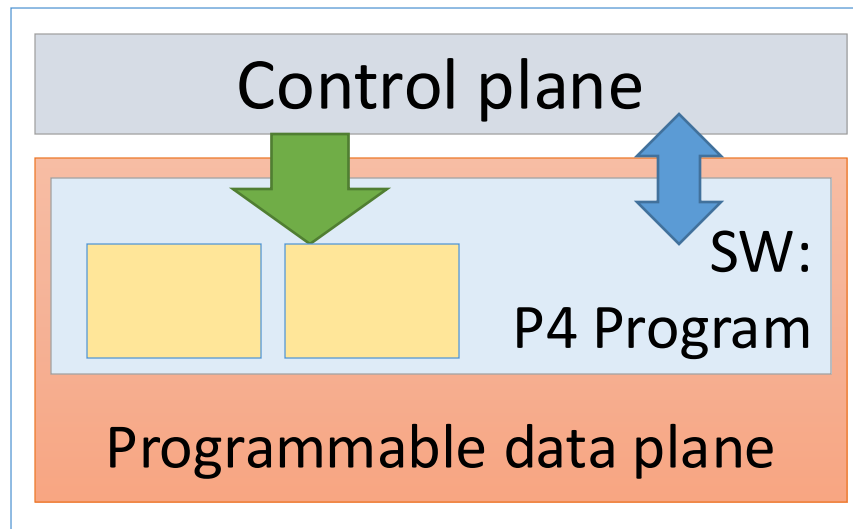
Software-Defined Networking



The P4 World



Not Just for Switches!



- Programmable switches
- FPGA switches
- Programmable network cards
- Software switches
- Hypervisor switches
- ...

Result:

- The flexibility of a SW program (with limitations)
- The performance of HW-based forwarding (depending on the target architecture)

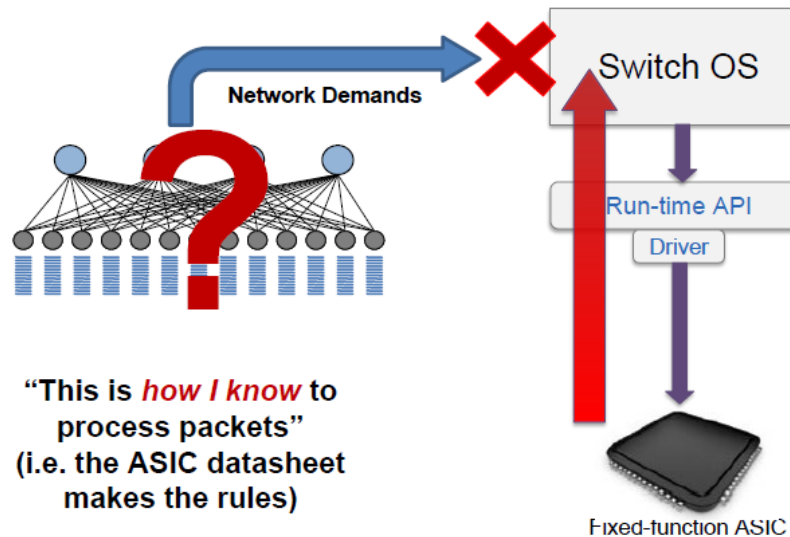
Programmable Network Devices

- PISA: Flexible Match+Action ASICs
 - Intel Flexpipe, Cisco Doppler, Cavium (Xpliant), Barefoot Tofino, ...
- NPU
 - EZchip, Netronome, ...
- CPU
 - Open vSwitch, eBPF, DPDK, VPP...
- FPGA
 - Xilinx, Altera, ...
- These devices let us tell them how to process packets

Intent-Based Programming

- Traditional network nodes expose their capabilities => not very programmable
- OpenFlow allows programmability, but still exposes a lot of internal details
- P4 programs using the desired behavior

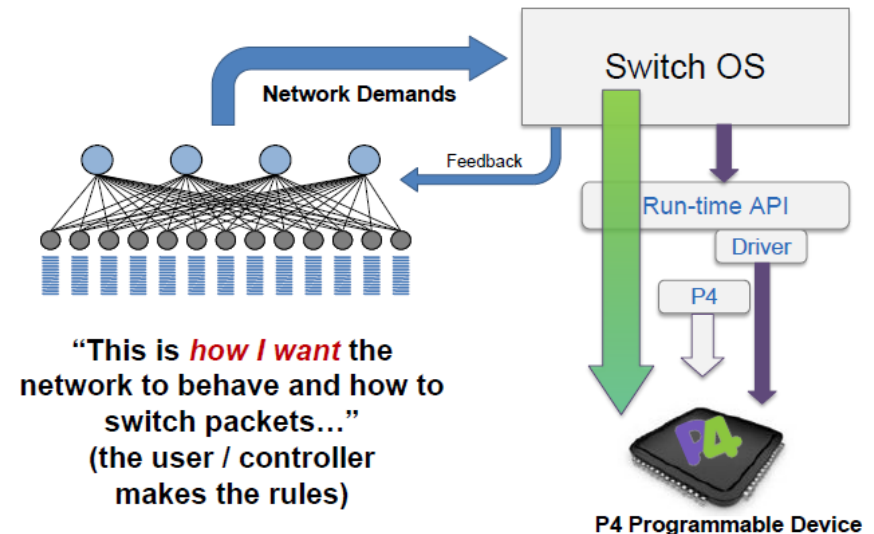
Status Quo: Bottom-up design



Copyright © 2017 – P4.org

3

A Better Approach: Top-down design



Copyright © 2017 – P4.org

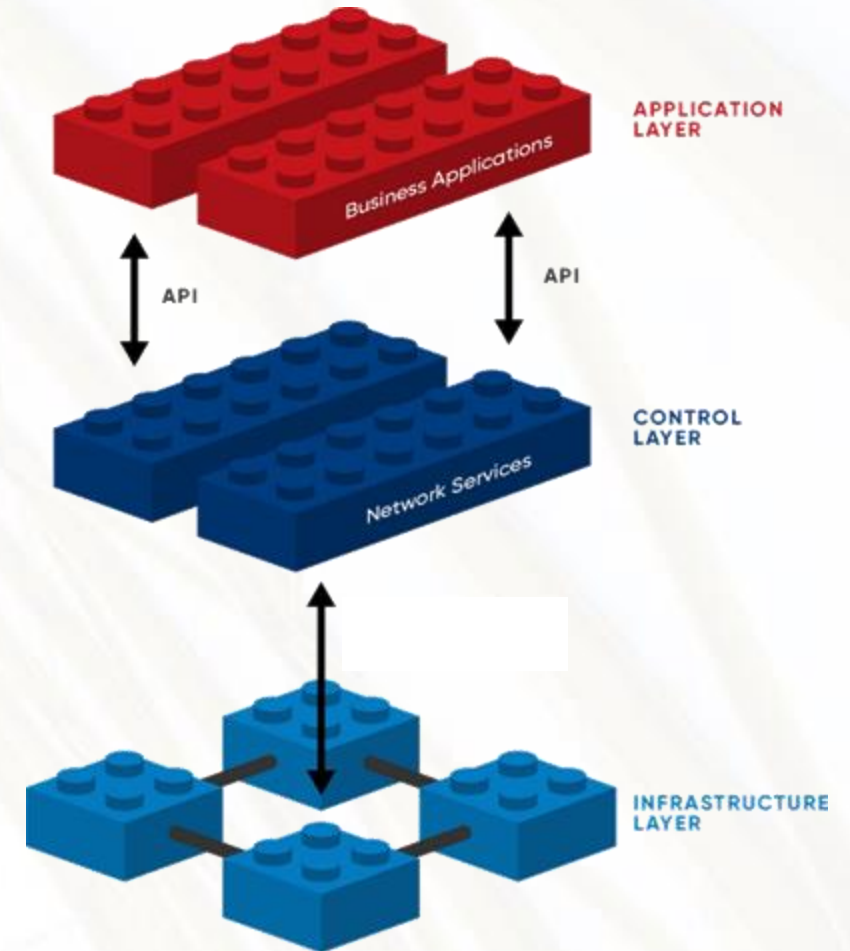
4



P4

Contents

- Why P4?
- What is P4?
- Standards & History
- Use Cases
- P4 Overview
- P4 Language Elements
- Discussion



P4 Community

- <http://github.com/p4lang>
- <http://p4.org>
- P4 Language Consortium
 - Mailing lists
 - Workshops
 - P4 developer days
 - Now a Project under ONF and the Linux Foundation
- Academic papers (SIGCOMM, SOSR)



*Carriers, cloud operators, chip vendors,
networking, universities, start-ups*

P4.org Consortium



P4 History

- May 2013: Initial idea and the name “P4”
- July 2014: First paper (SIGCOMM ACR)
- Aug 2014: First P4₁₄ Draft Specification (v0.9.8)
- Sep 2014: P4₁₄ Specification released (v1.0.0)
- May 2017: P4₁₄ v1.0.4
- Apr 2016: P4₁₆ – first commits
- May 2017: P4₁₆ Specification released
- Oct 2019: P4₁₆ Latest Specification released
- . . .
- Official Spelling P4_16 or P4₁₆

P4₁₆

- Most recent revision of P4
- Similar to C; strongly typed
- Currently in draft form
- Spec: <https://p4.org/p4-spec/docs/P4-16-v1.2.0.pdf>
- Reference compiler implementation (Apache 2 license): <http://github.com/p4lang/p4c>



Available Software Tools

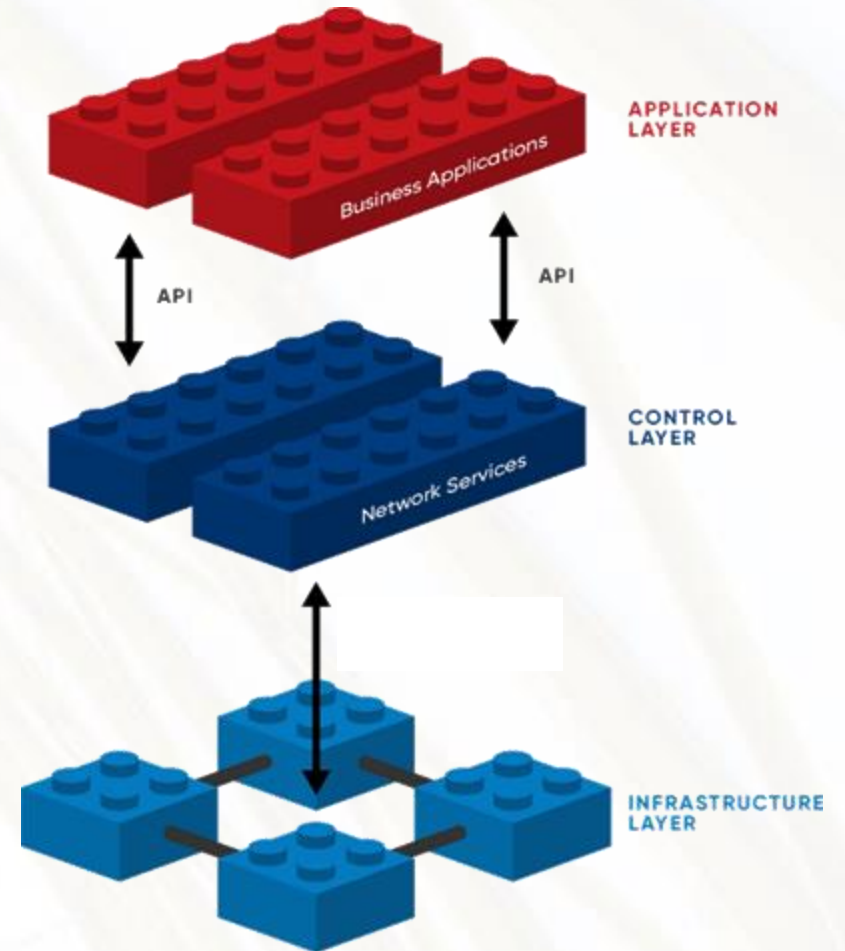
- Compilers for various back-ends
 - Netronome chip, Barefoot chip, BMv2, eBPF, Xilinx FPGA (open-source and proprietary)
- Multiple control-plane implementations
 - SAI, OpenFlow
- Simulators
- Testing tools
- Sample P4 programs
- Tutorials



P4

Contents

- Why P4?
- What is P4?
- Standards & History
- Use Cases
- P4 Overview
- P4 Language Elements
- Example
- Discussion



Use Only What You Need

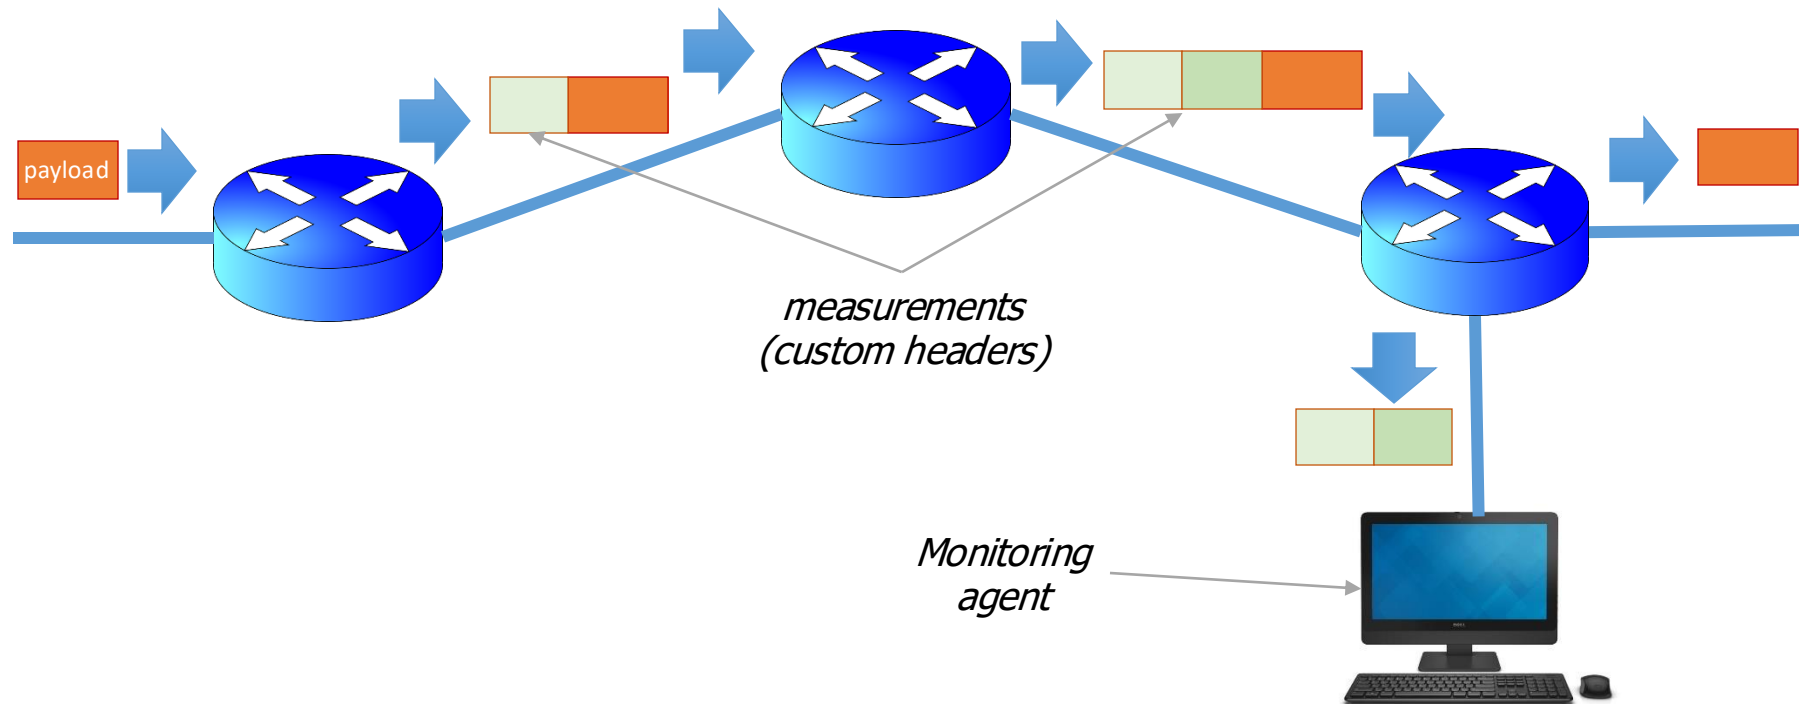
- IETF has issued thousands of RFCs
- Switch RAM and CPU is very expensive
- Network operators can *remove* protocols
- Simpler troubleshooting



What Can You do With P4?

- Layer 4 Load Balancer – SilkRoad[1]
 - Low Latency Congestion Control – NDP[2]
 - Fast In-Network cache for key-value stores – NetCache[3]
 - In-band Network Telemetry – INT[4]
 - Consensus at network speed – NetPaxos[5]
 - ... and much more
-
- [1] Miao, Rui, et al. "SilkRoad: Making Stateful Layer-4 Load Balancing Fast and Cheap Using Switching ASICs." SIGCOMM, 2017.
 - [2] Handley, Mark, et al. "Re-architecting datacenter networks and stacks for low latency and high performance." SIGCOMM, 2017.
 - [3] Xin Jin et al. "NetCache: Balancing Key-Value Stores with Fast In-Network Caching." SOSP 2017
 - [4] Kim, Changhoon, et al. "In-band network telemetry via programmable dataplanes." SIGCOMM. 2015.
 - [5] Dang, Huynh Tu, et al. "NetPaxos: Consensus at network speed." SIGCOMM, 2015.

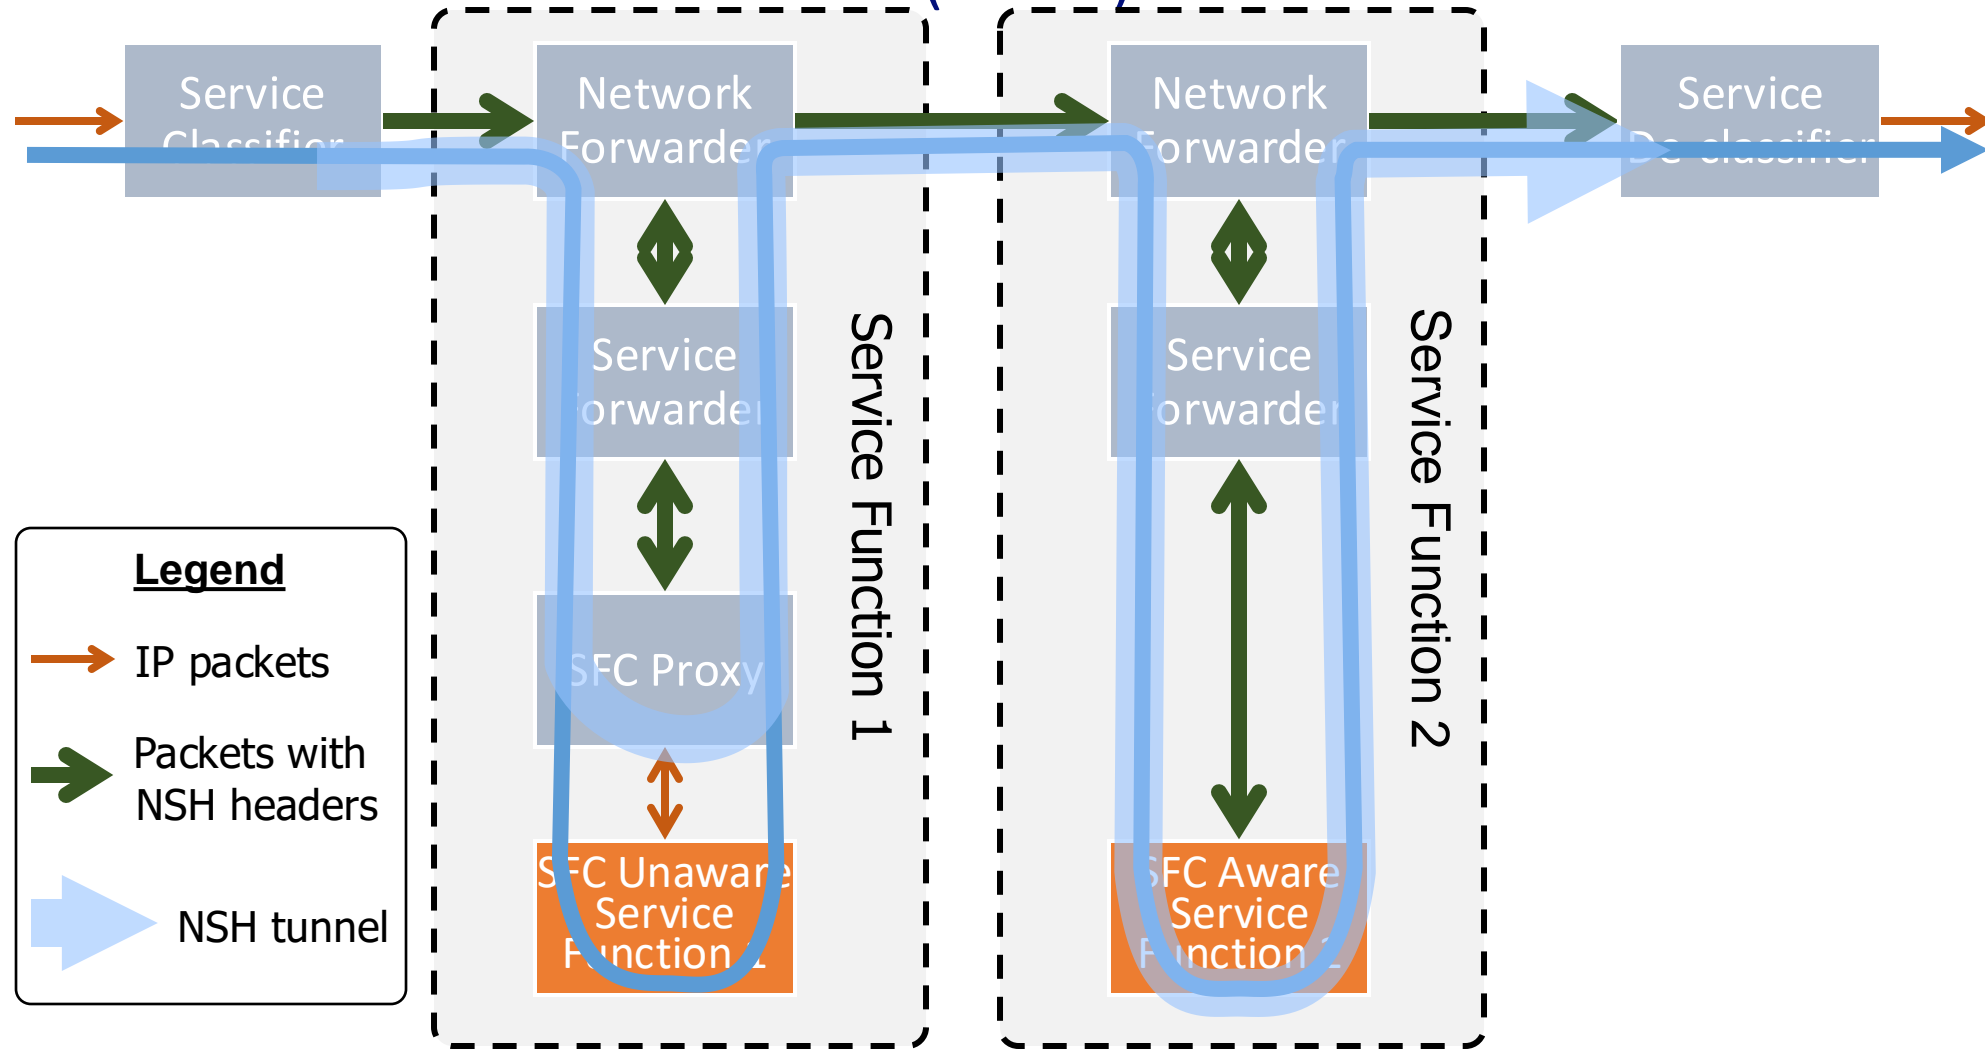
Network Monitoring



In-Band Network Telemetry (INT)
**Improving Network Monitoring and Management
with Programmable Data Planes**



By Mukesh Hira & LJ Wobker

IETF Service Function Chaining: Network Services Header (NSH)



Network = Software

- Use **software** engineering principles and tools
- Upgrade your network at any time
- Protocols = intellectual property



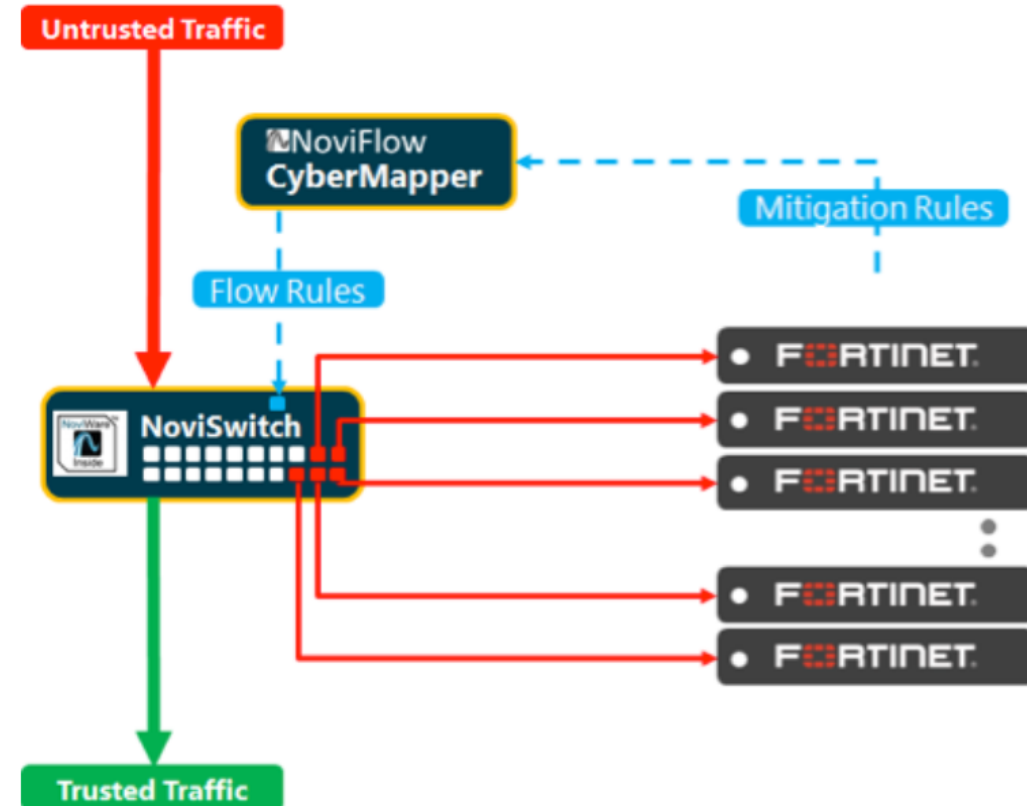
Protocols = Programs

- Implement (new) protocols
 - VxLAN: 175 lines of code
 - NVGRE: 183 lines of code
- Low overhead (high speed)
- Define your own packet processing policies
- Improved signaling, monitoring, and troubleshooting
- Change functionality with software upgrades
- Use only what you need



Scaling Network Security Services

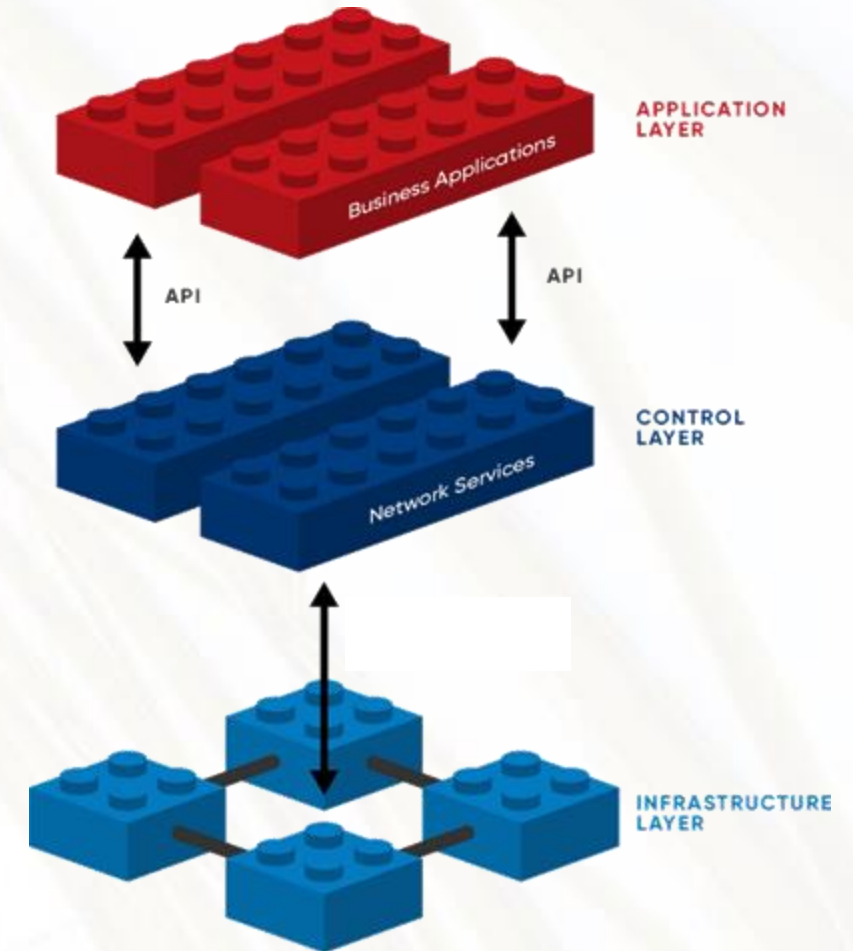
- Services
 - Threat Intelligence Gateway
 - Packet Broker
 - Filtering, mirroring, Port-pairing
 - Load Balancing
 - Sticky stateless load balancing
 - Proportional load balancing
 - Pool Scaling
- Networking
 - In-band Network Telemetry (INT)
 - Service Chaining
 - Segment Routing



P4

Contents

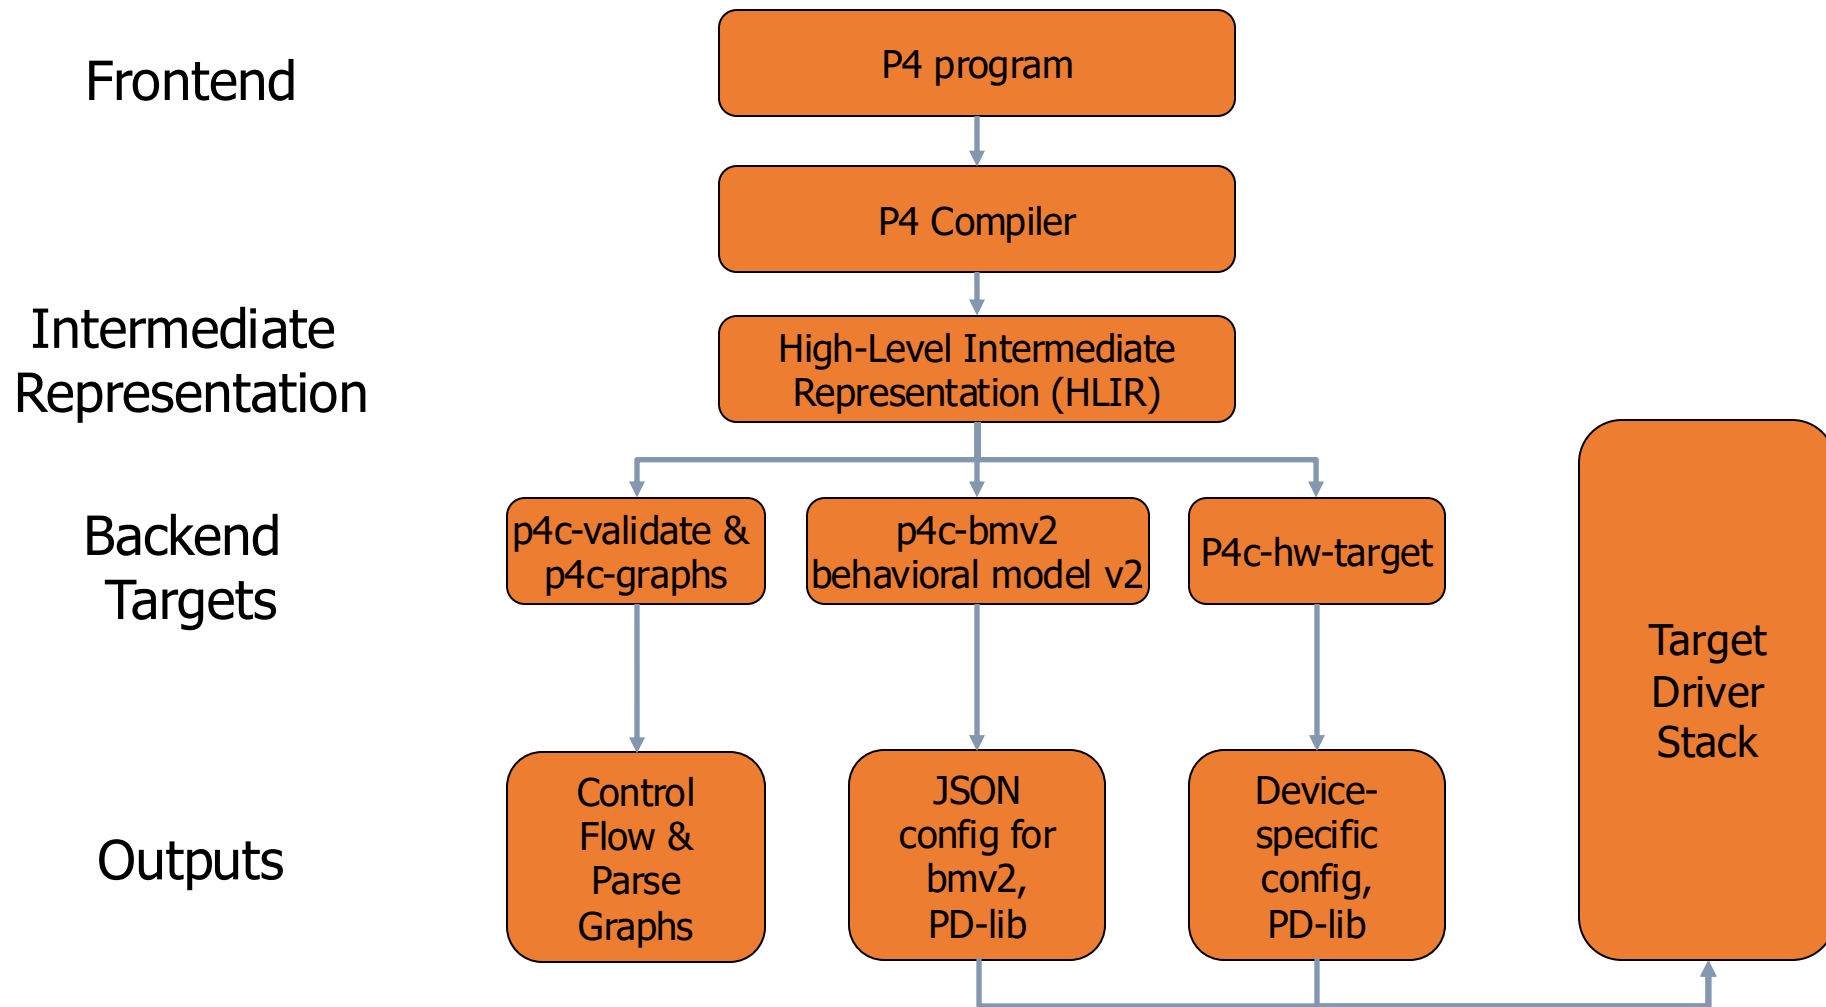
- Why P4?
- What is P4?
- Standards & History
- Use Cases
- P4 Overview
- P4 Language Elements
- Discussion



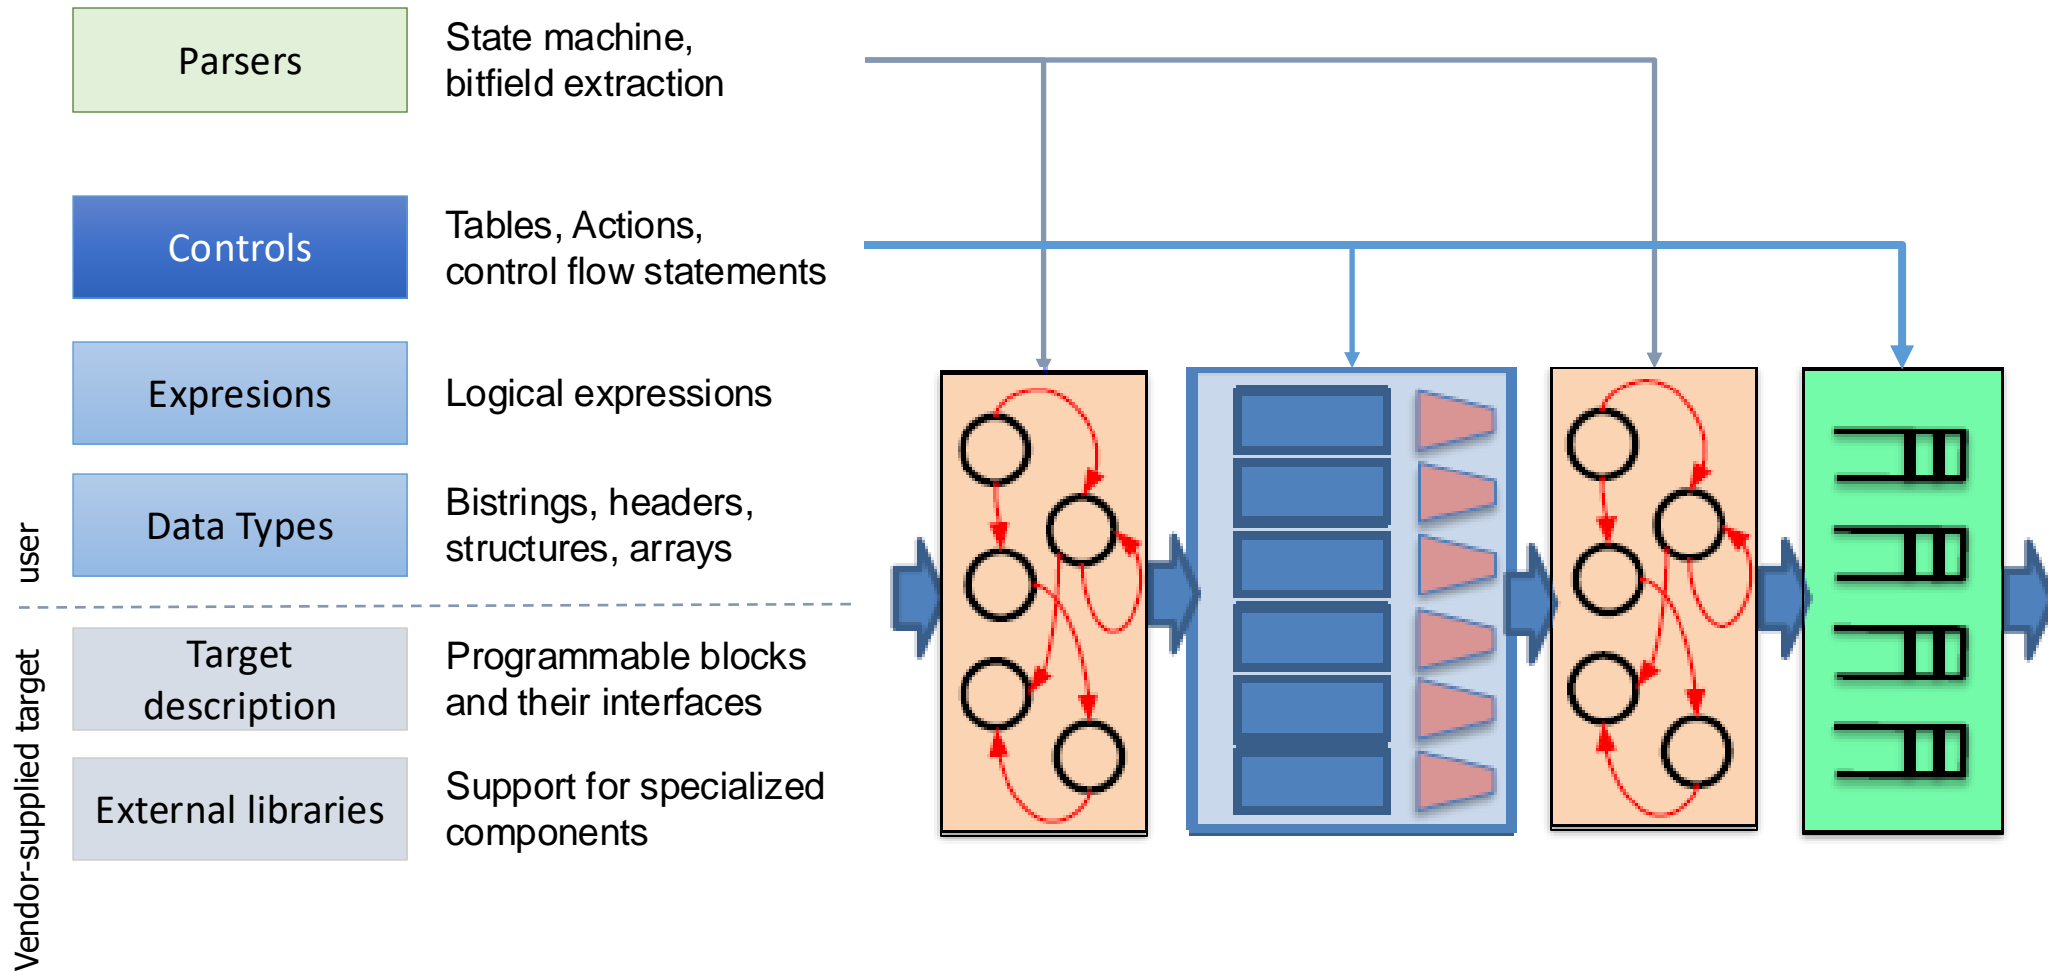
P4 - Protocol Independent Packet Processing

- Protocol independence
 - Configure packet parser. Not restricted to few protos
 - Define a set of typed match + action tables
- Target independence
 - Program without knowledge of switch details
 - Rely on compiler to configure the target switch
- Reconfigurability
 - Change parsing and processing in the field
- Programmable Network Devices
 - PISA: Protocol Independent Switch Architecture
 - ASIC, NPU, CPU, FPGA


P4 Modular Compilation



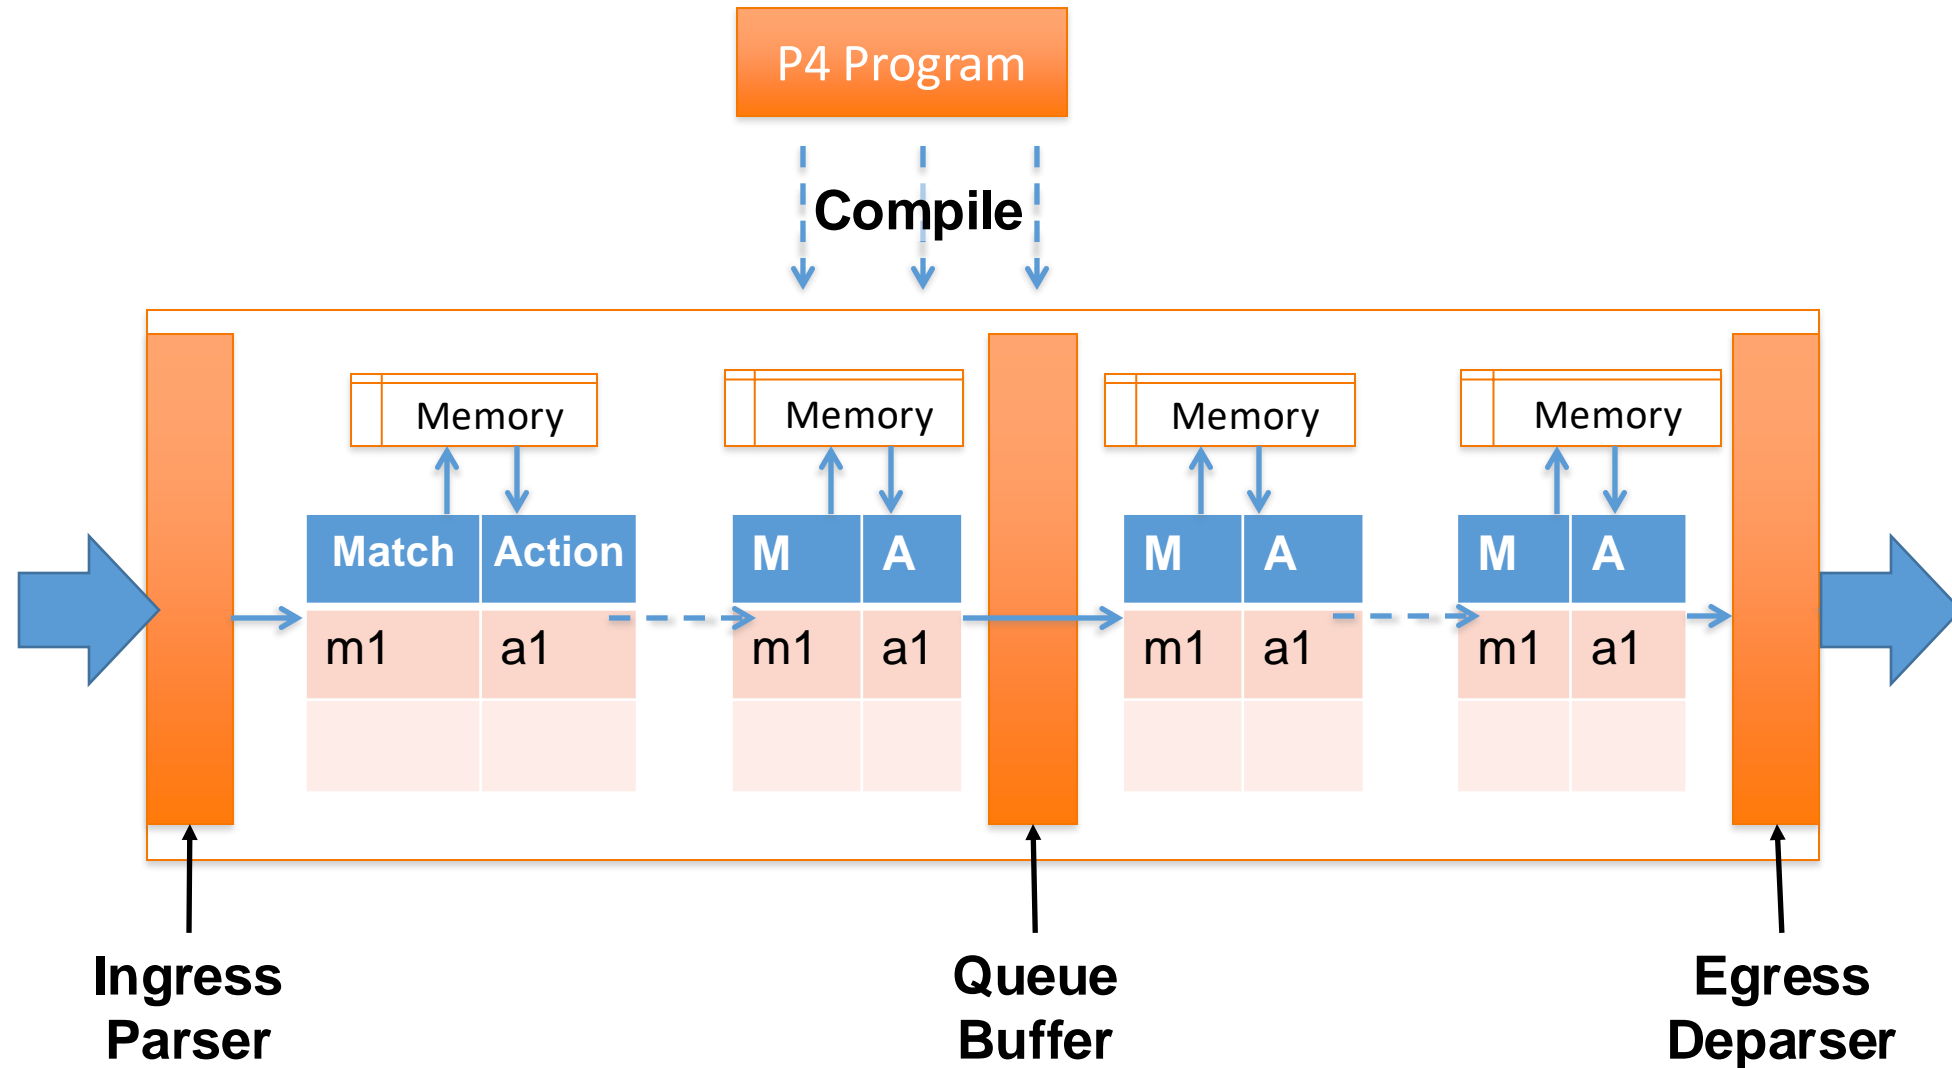
P4₁₆ Language Elements



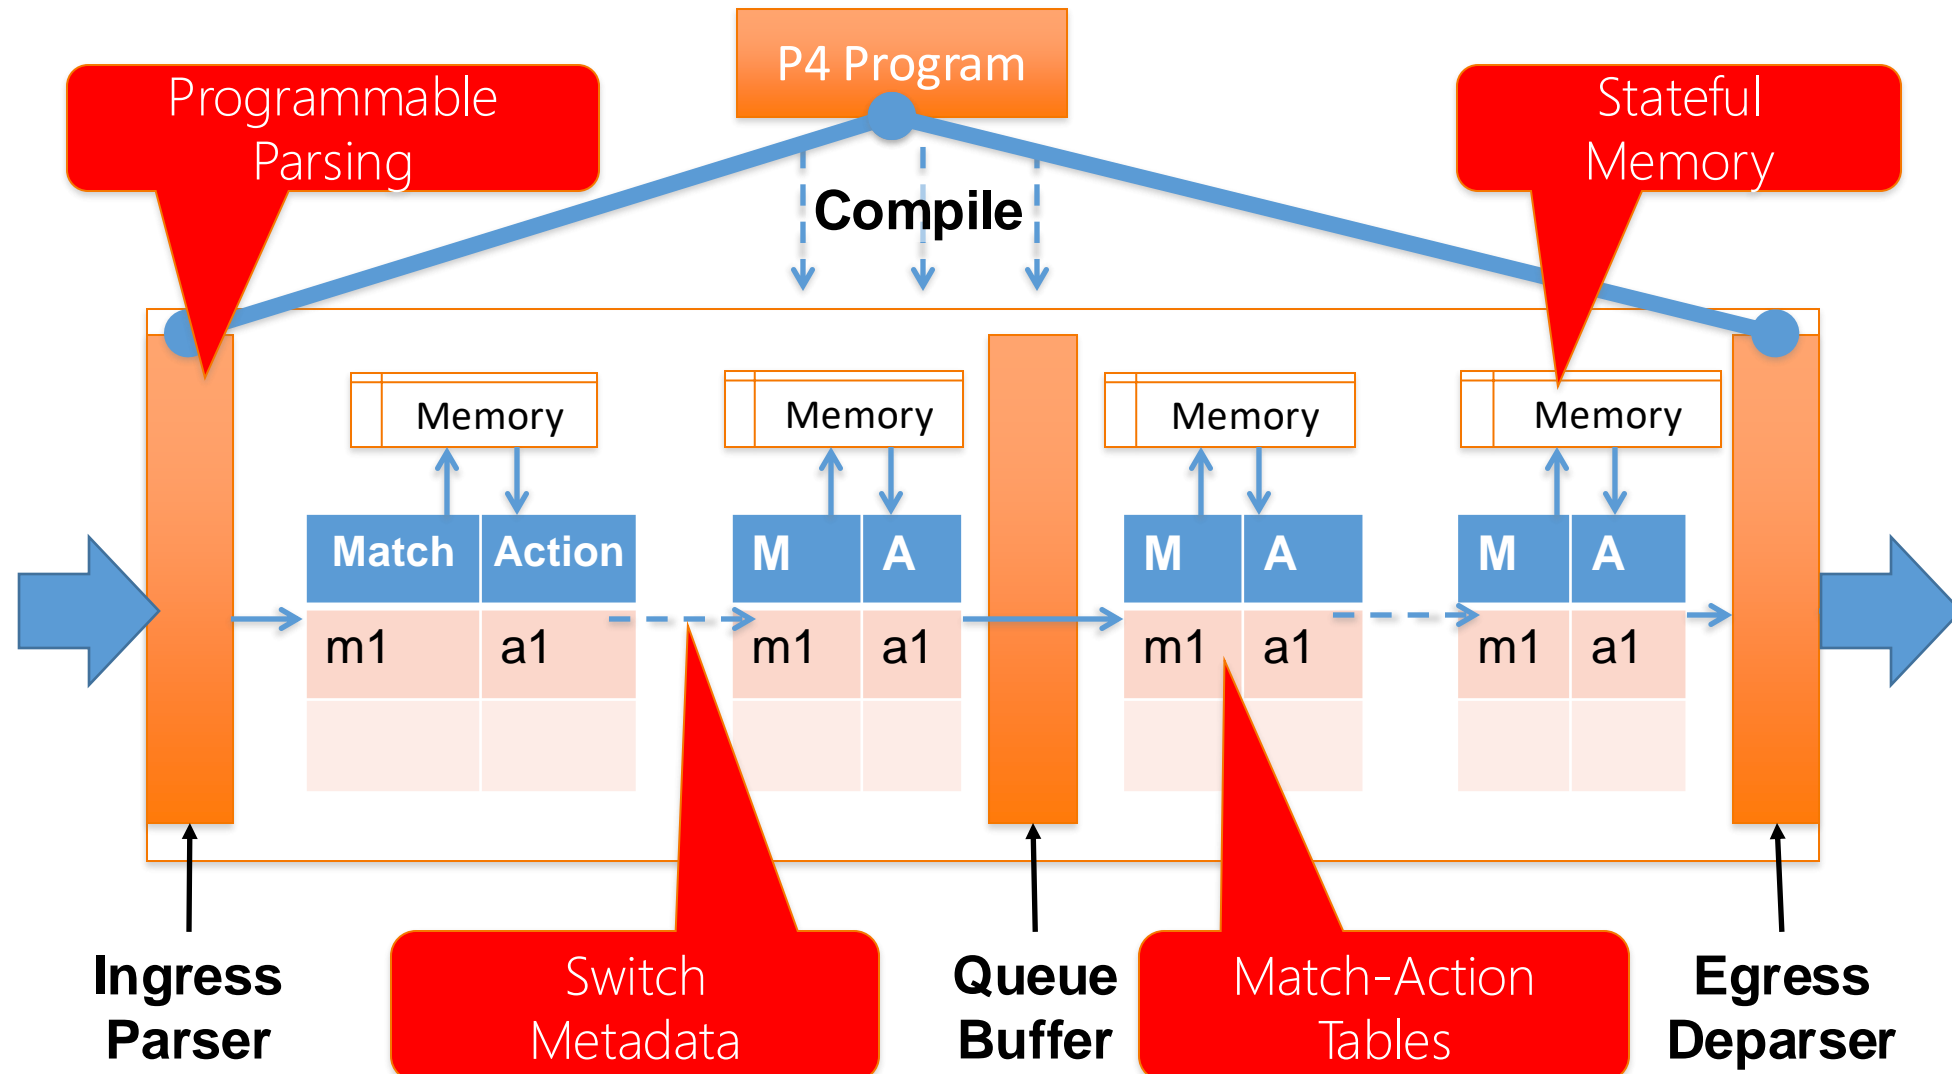
P4₁₆ Data Plane Model



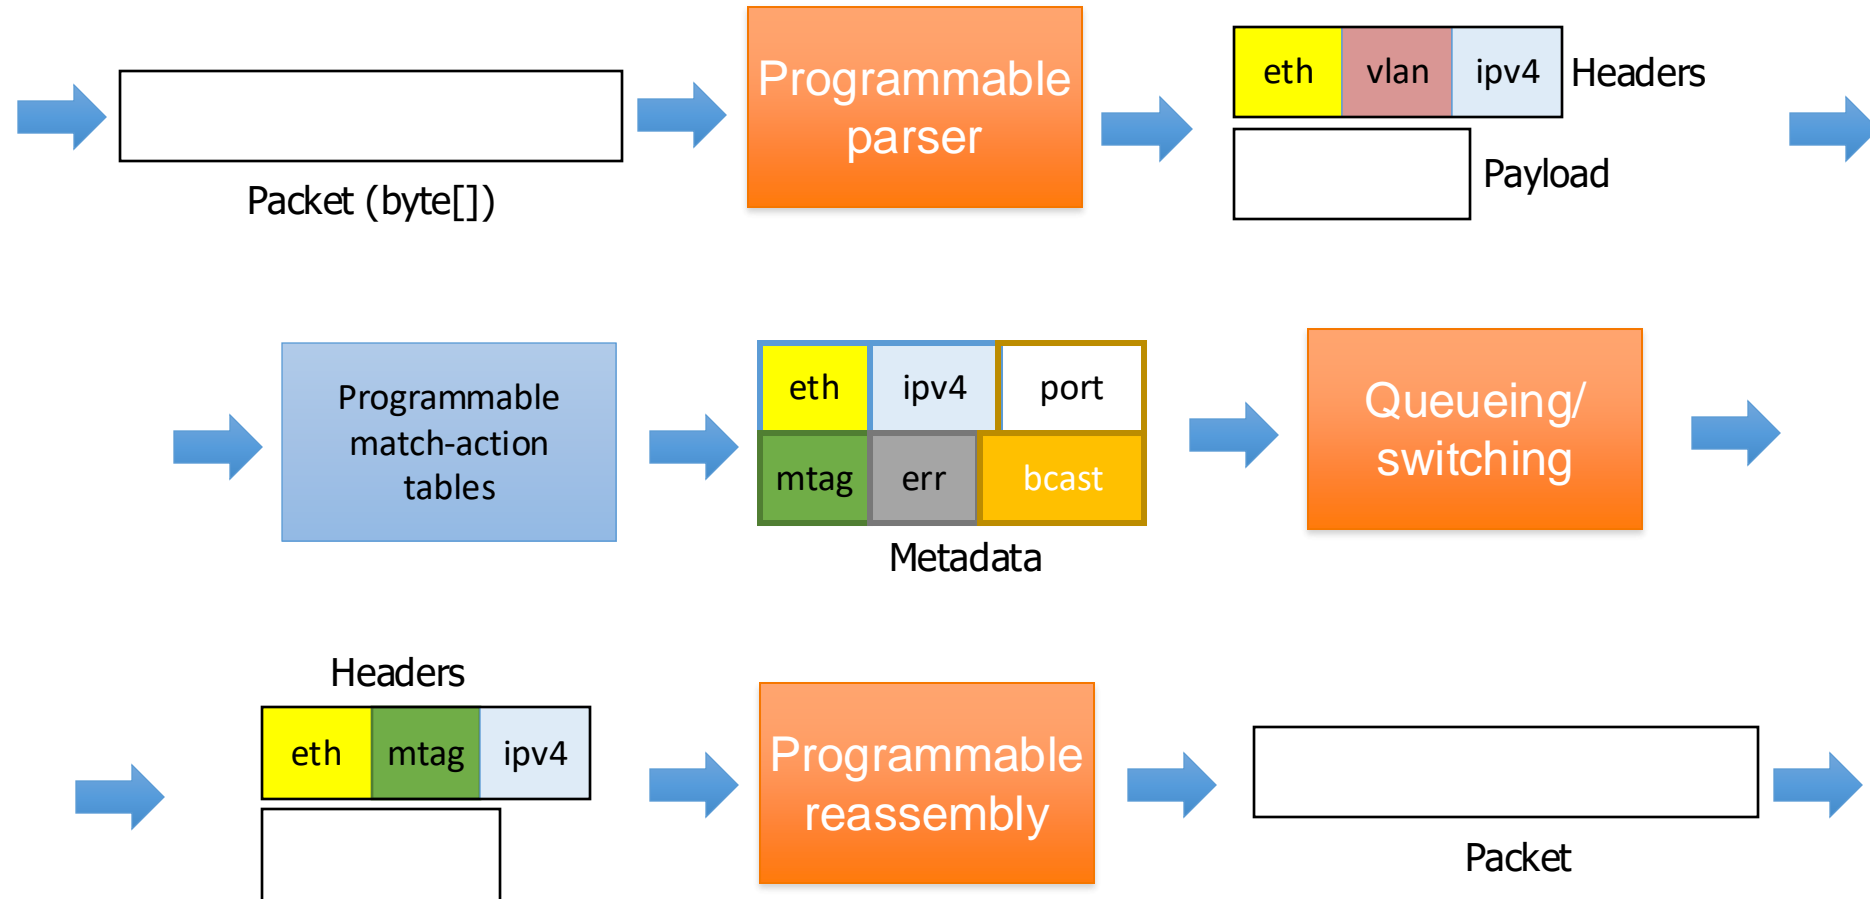
Programmable Switches: Capabilities



Programmable Switches: Capabilities

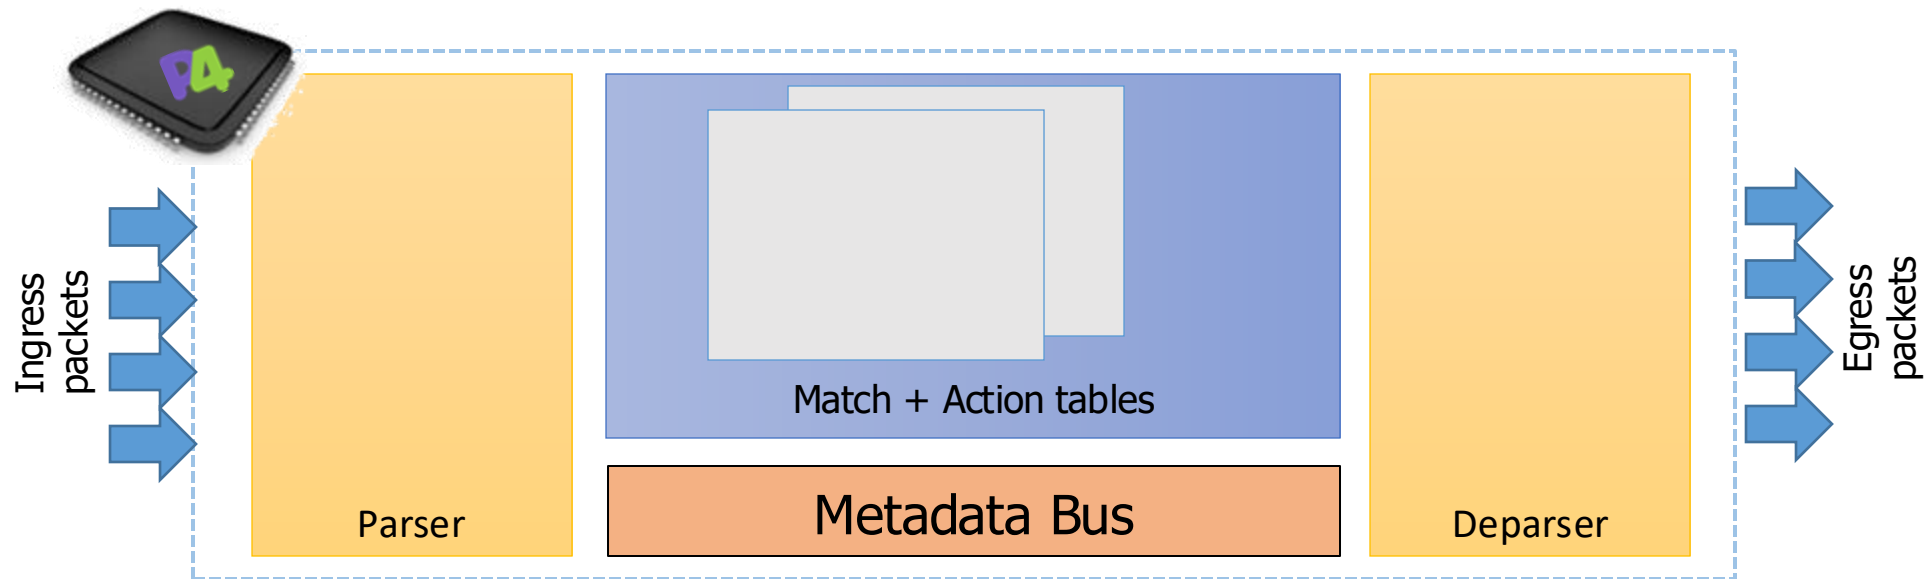


Example: Packet Processing Pipeline

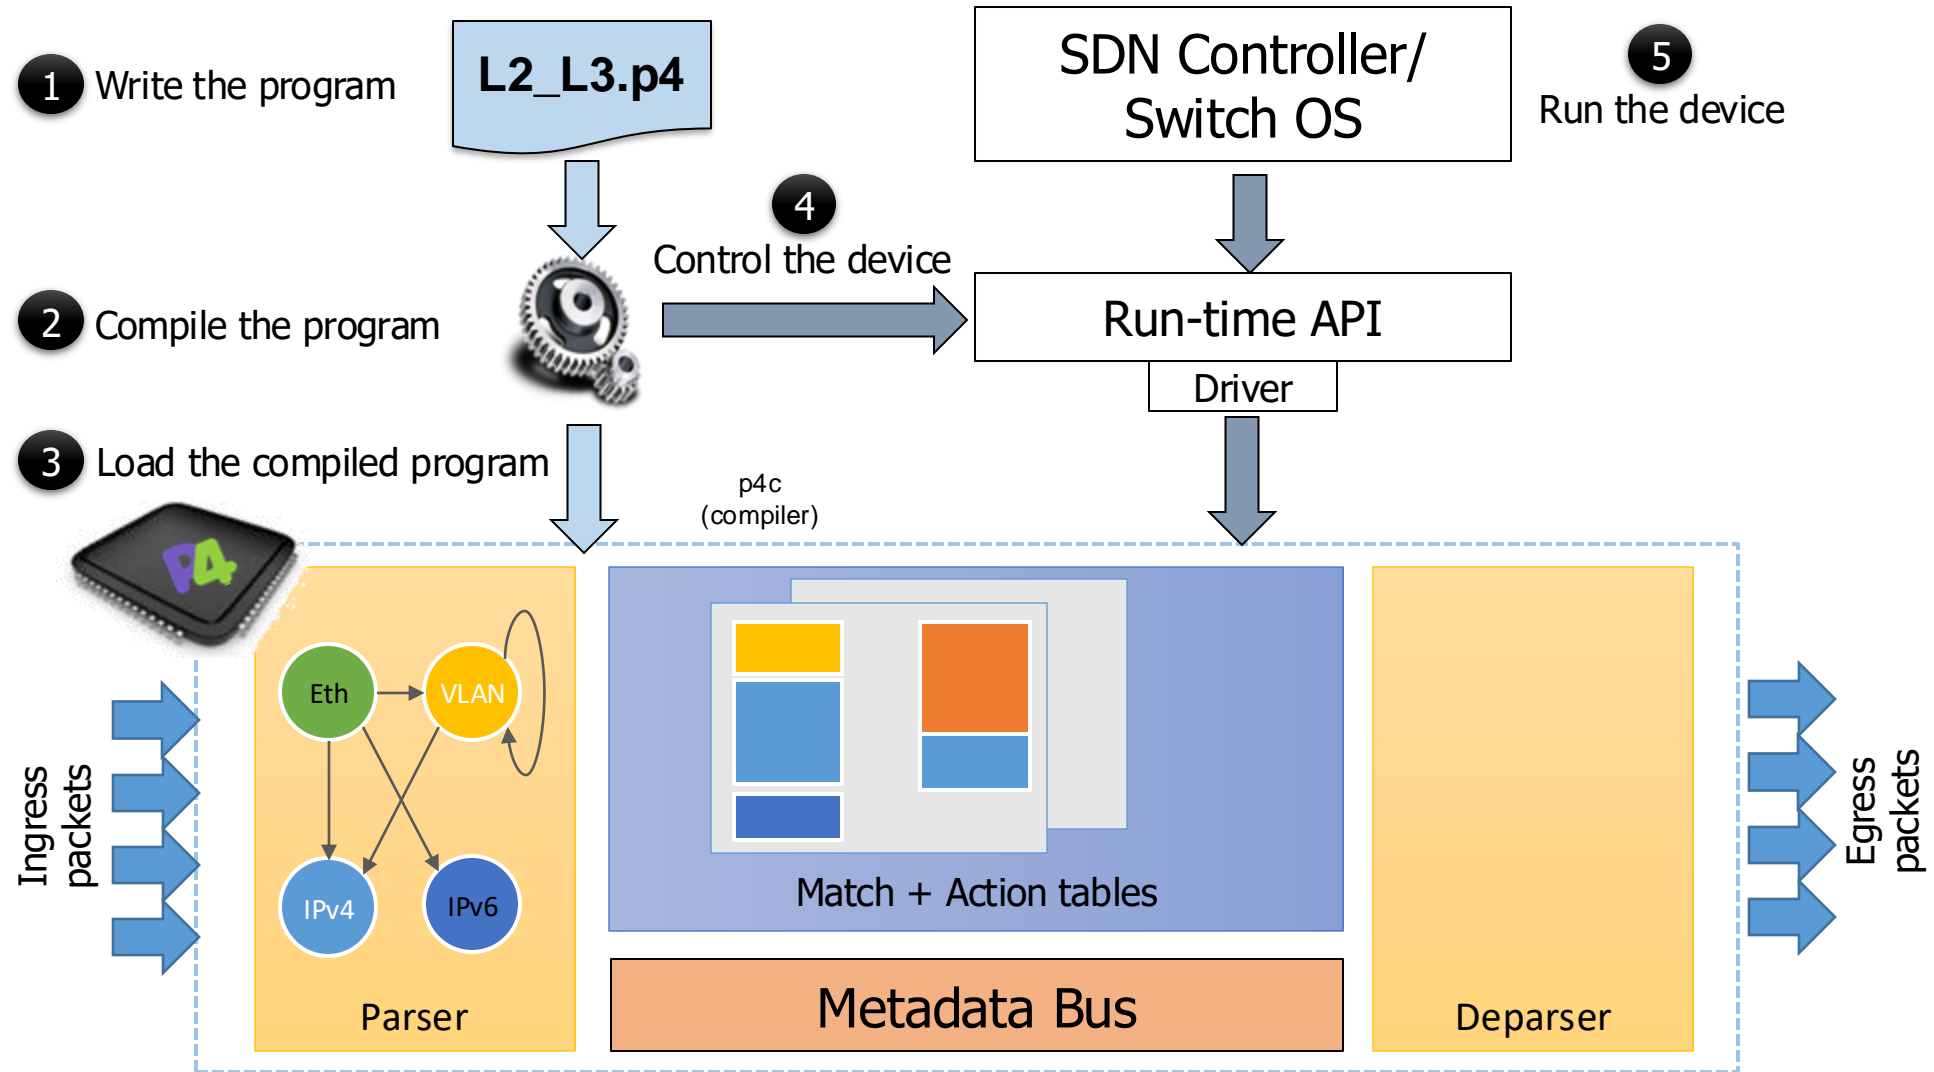


P4-Based Workflow

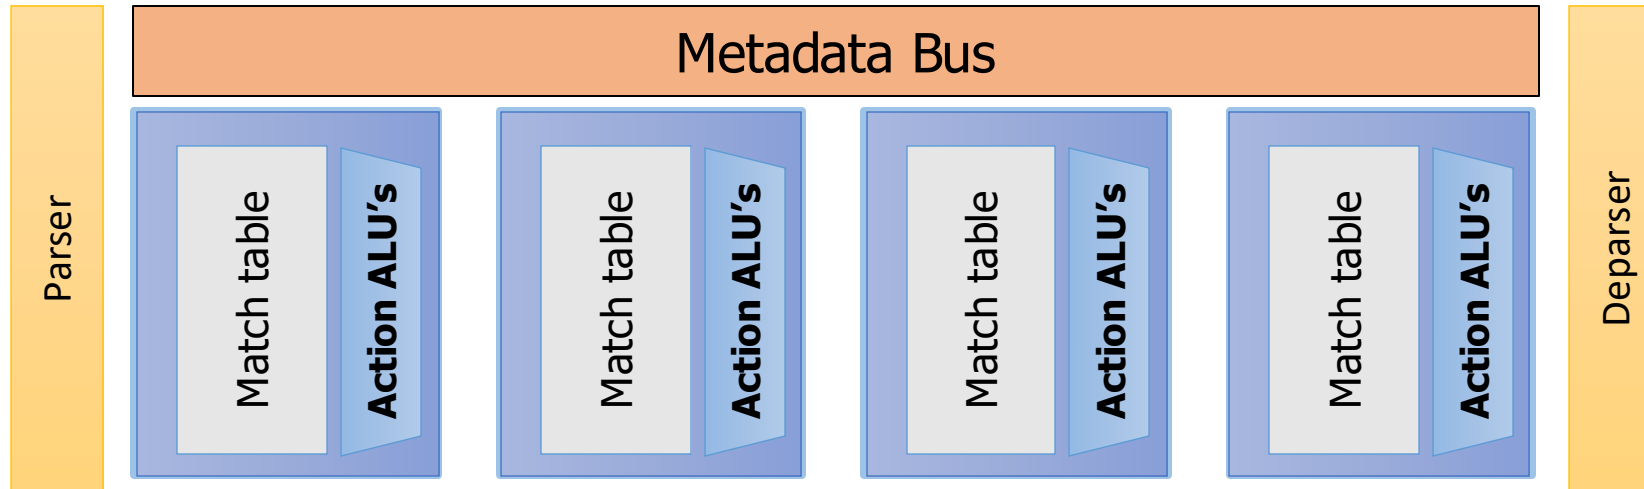
- Device is not programmed yet
 - Parser, tables and deparser are empty
 - Does not know about any packet formats or protocols



P4-Based Workflow



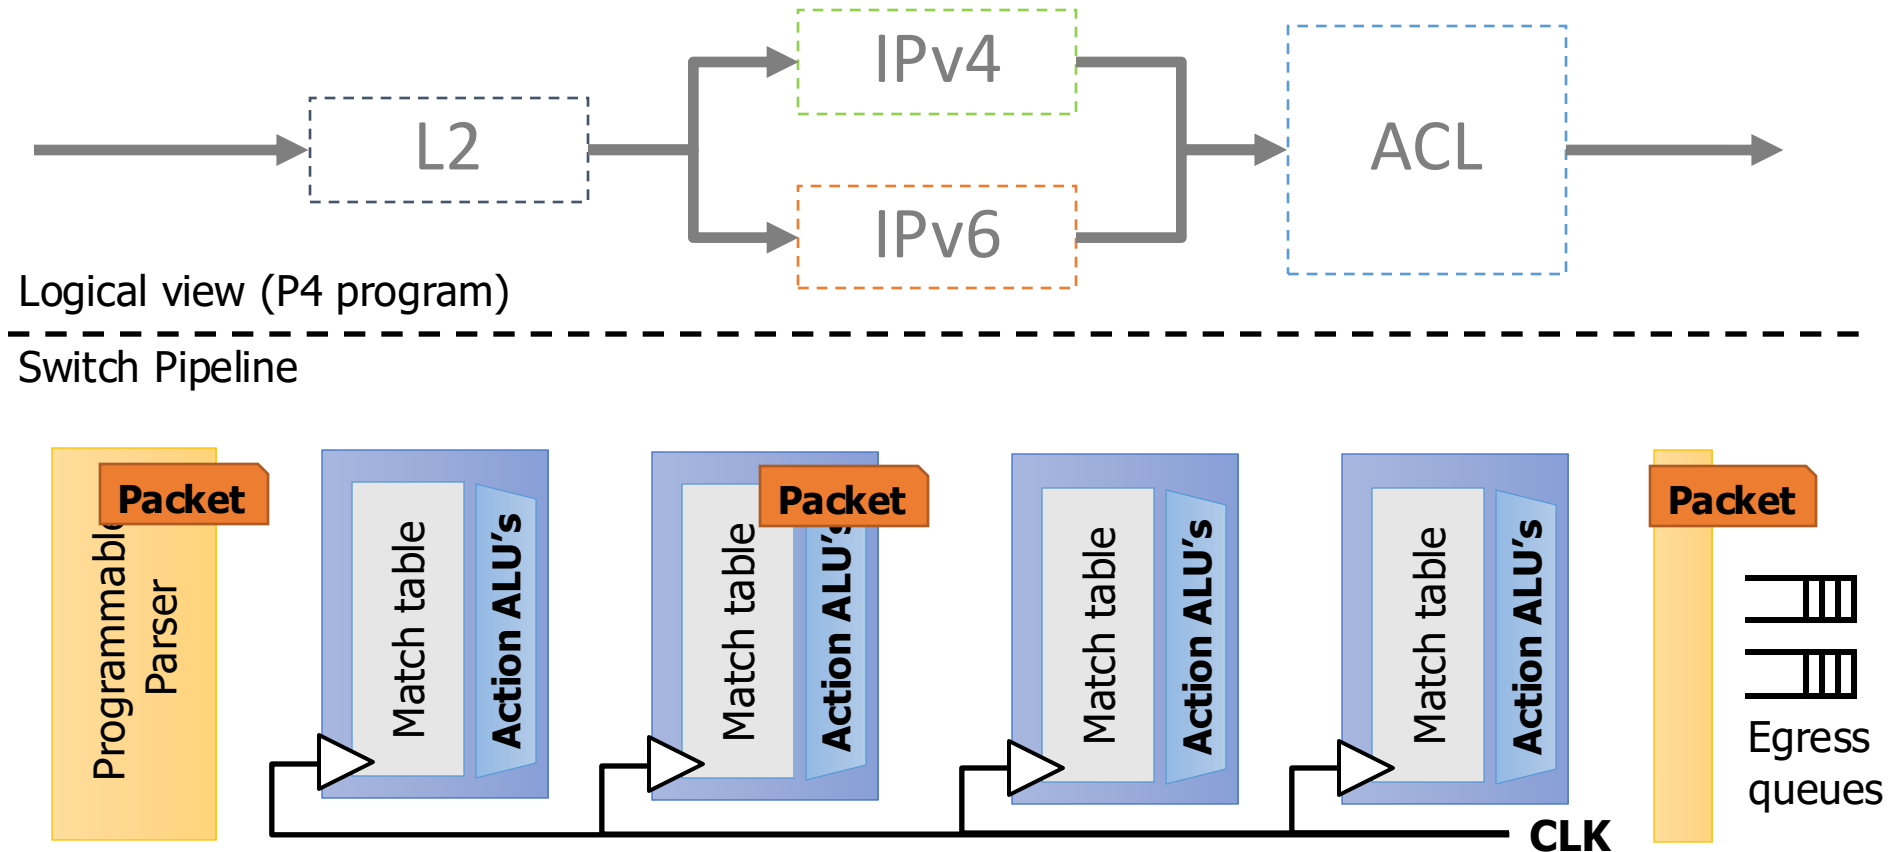
The anatomy of a basic pipeline



- Parser
 - Converts packet data into a metadata (Parsed representation)
- Match + Action Tables
 - Operate on metadata
- Deparser
 - Converts metadata back into a serialized packet
- Metadata Bus
 - Carries the information within the pipeline

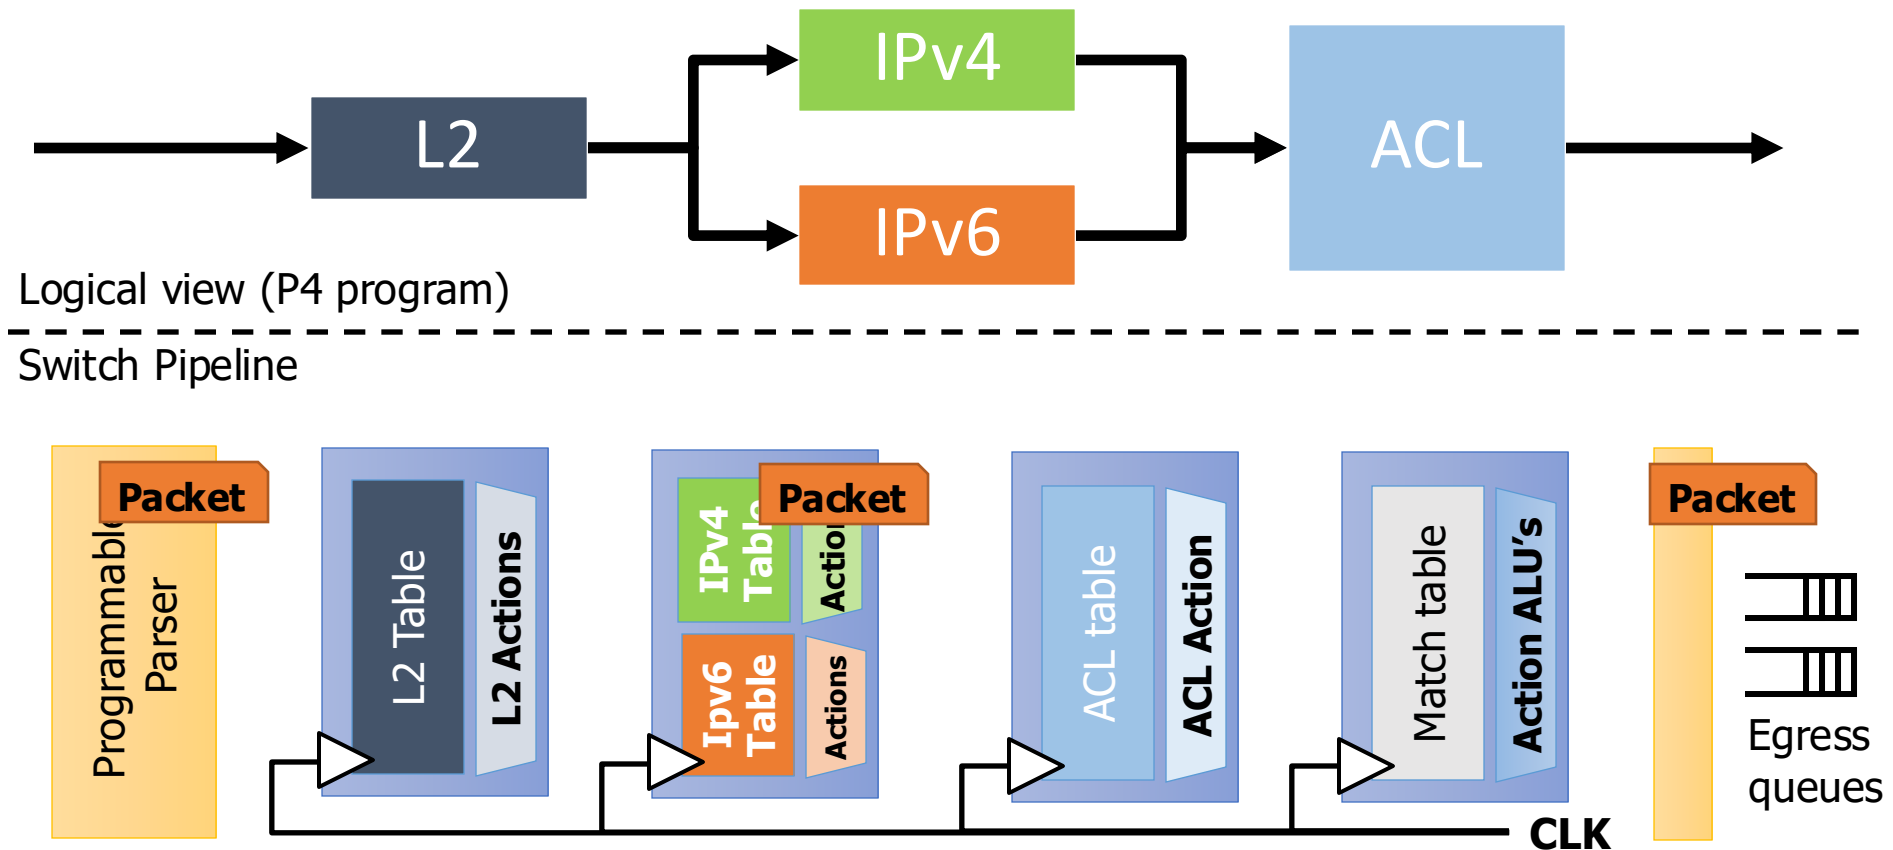
Mapping a P4 Program into a Pipeline

- No Program has been downloaded yet



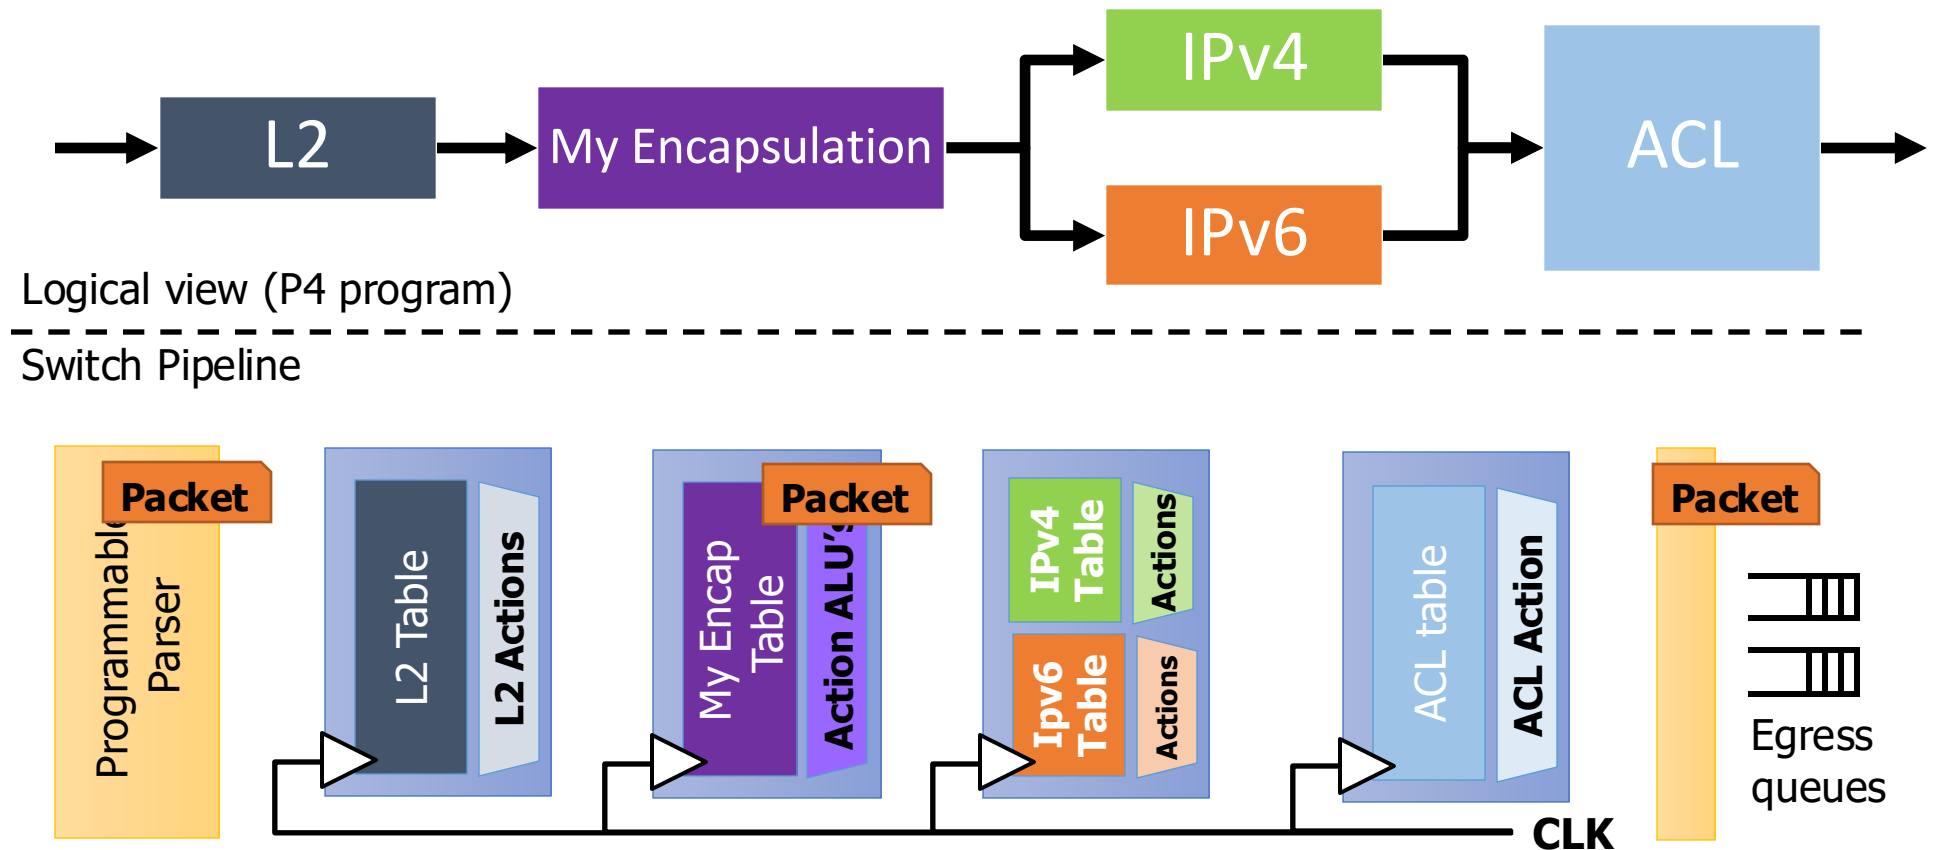
Mapping a P4 Program into a Pipeline

- Program has been compiled and downloaded

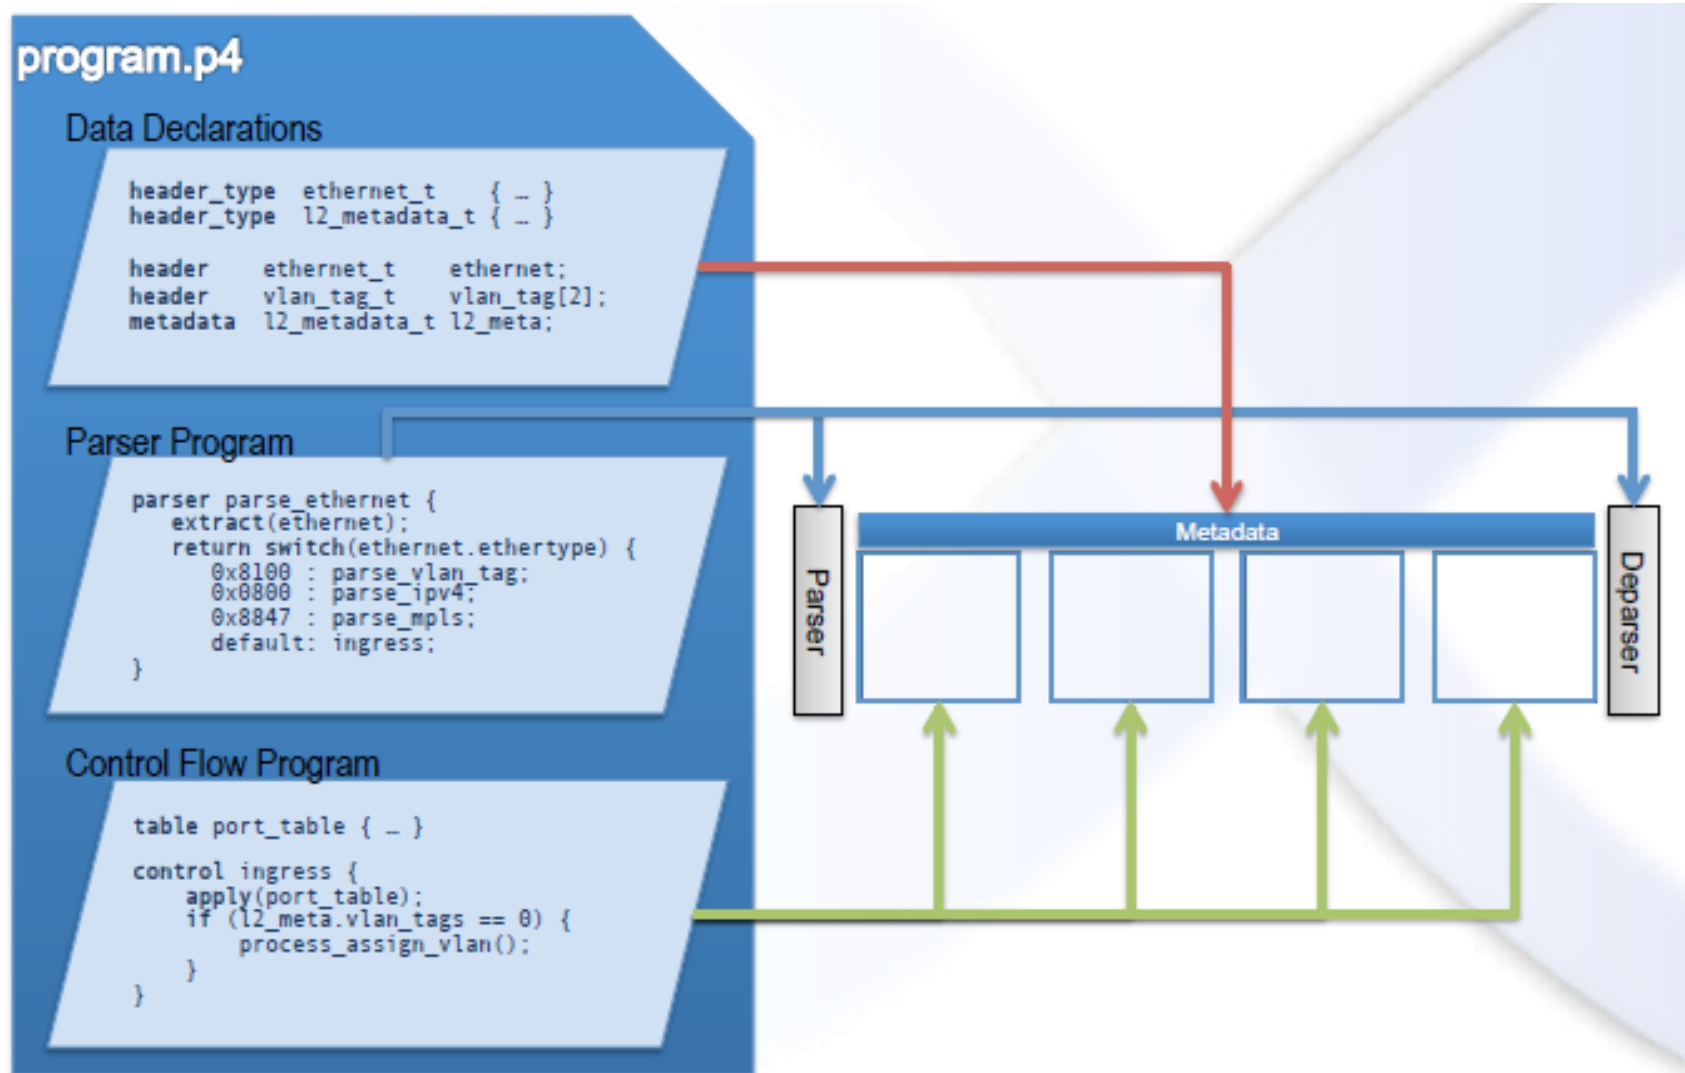


Mapping a P4 Program into a Pipeline

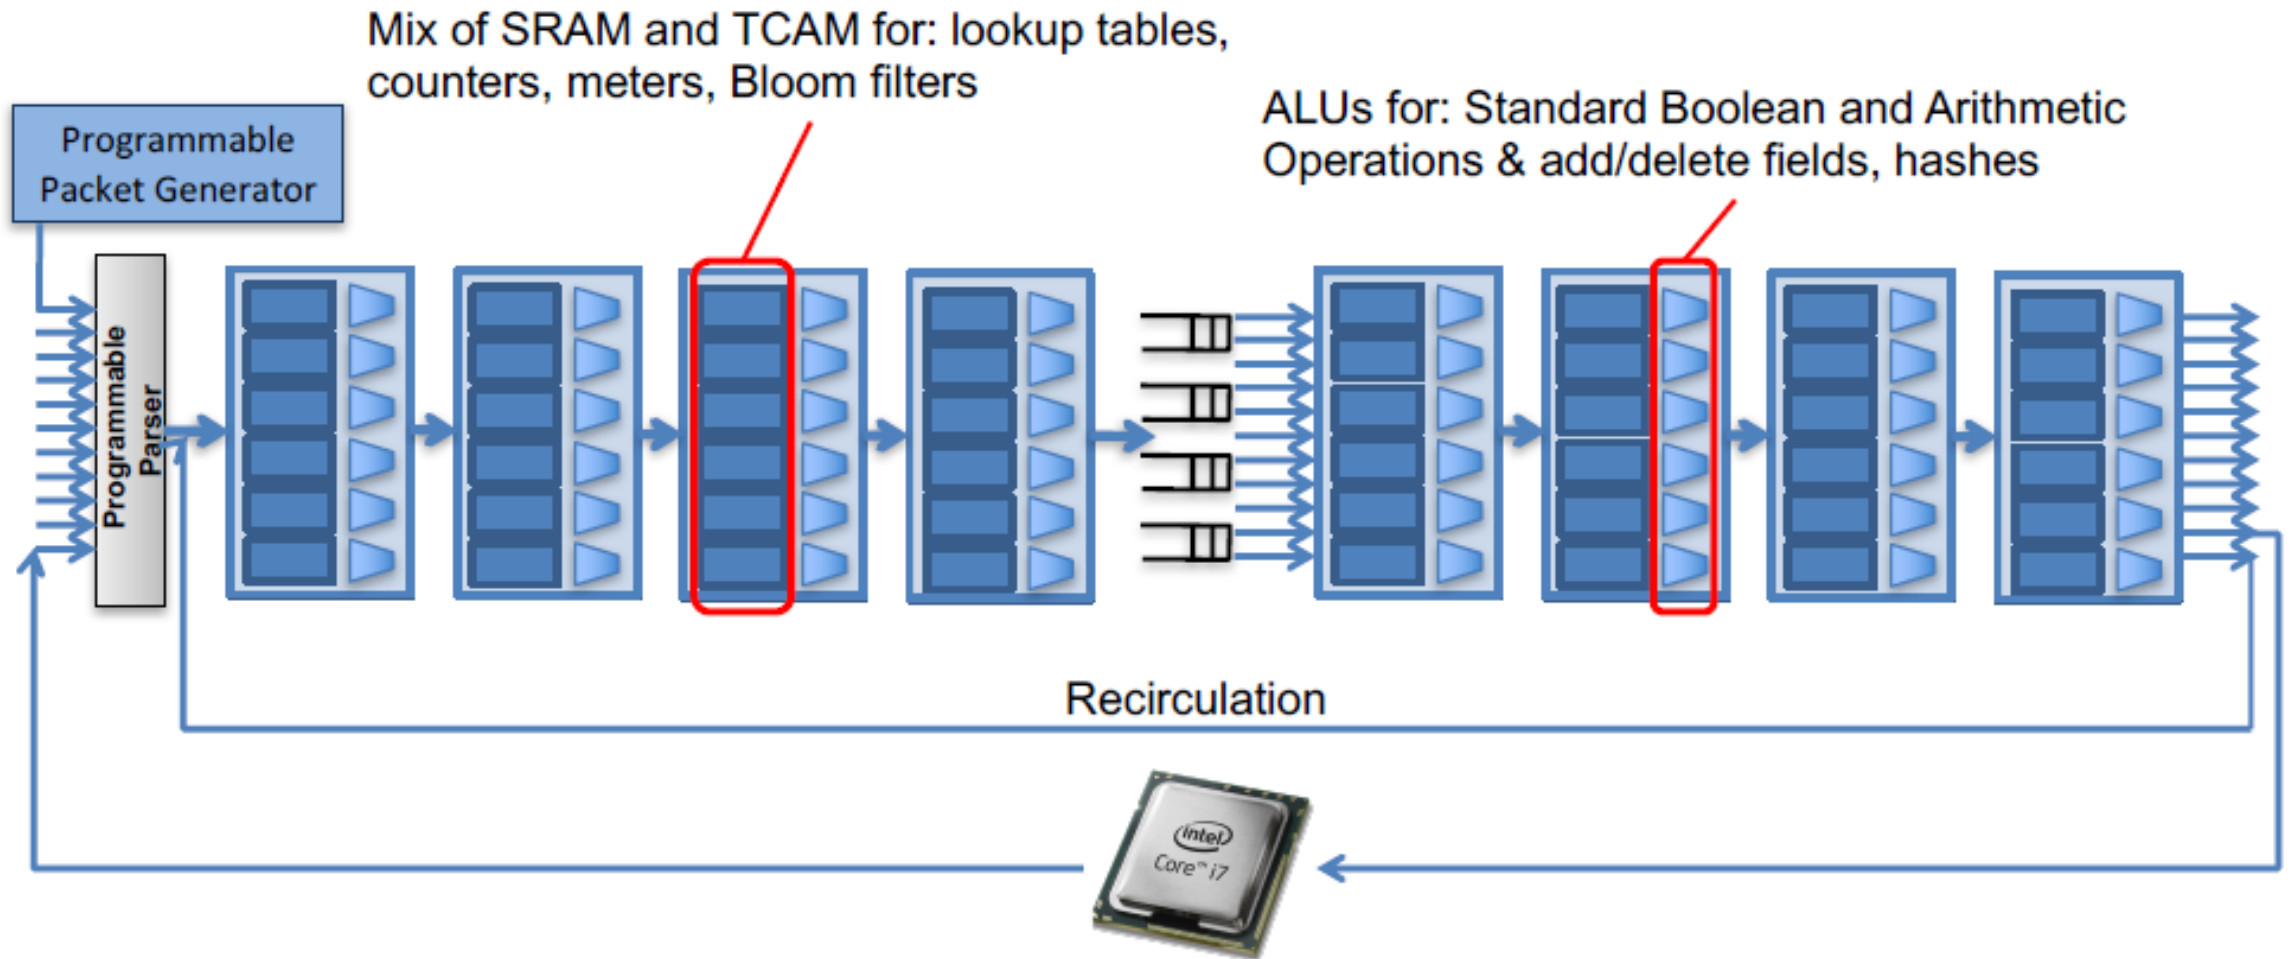
- Re-Program in the field



P4 Program Sections



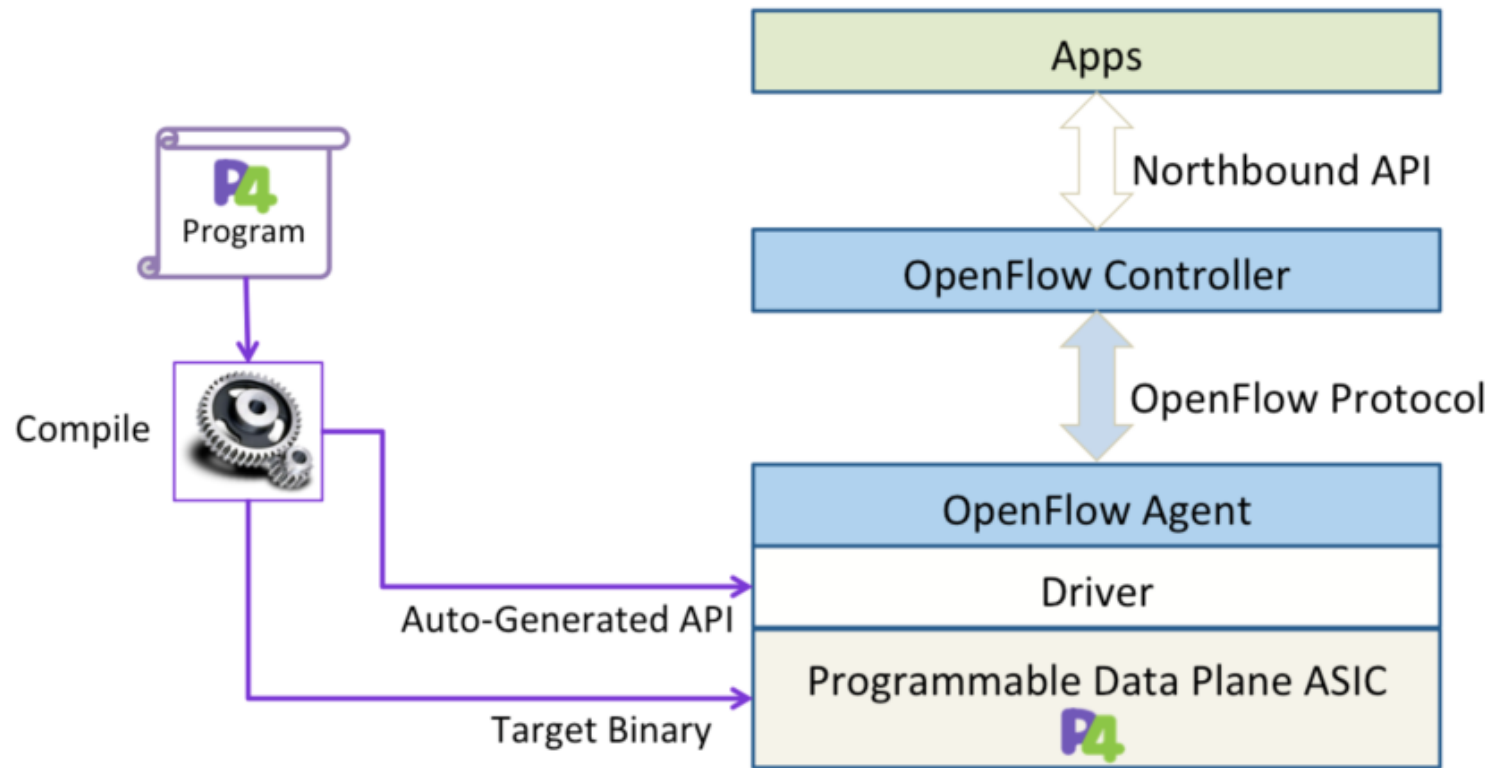
PISA: Protocol Independent Switch Architecture



P4 & OpenFlow



P4 & OpenFlow

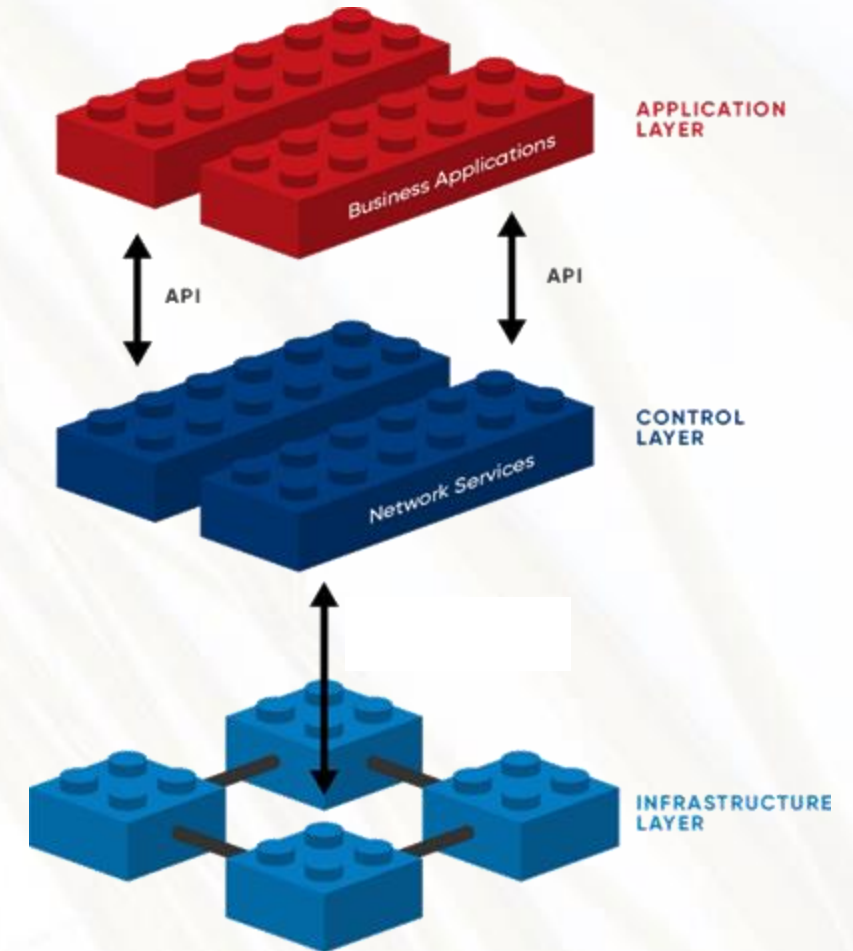


Copyright © 2016 P4 Language Consortium.

P4

Contents

- Why P4?
- What is P4?
- Standards & History
- Use Cases
- P4 Overview
- P4 Language Elements
- Discussion



Why P4₁₆

- Clearly defined semantics
 - You can describe what your data plane program is doing
- Expressive
 - Supports a wide range of architectures
- High-level, Target-independent
 - Uses conventional constructs
 - Compiler manages the resources and deals with the hardware
- Type-safe
 - Enforces good software design practices & eliminates “stupid” bugs
- Agility
 - High-speed networking devices become as flexible as any software
- Insight
 - Freely mixing packet headers and intermediate results

P4 Language Components

- Data declarations
 - Packet Headers and Metadata
- Parser Programming
 - Parser Functions (Parser states)
 - Checksum Units
- Packet Flow Programming
 - Actions
 - Primitive and compound actions
 - Counters, Meters, Registers
 - Tables
 - Match keys & Attributes
 - Control Functions (Imperative Programs)
- No: pointers, loops, recursion, floating point!!

Limitations of P4₁₆

- The core P4 language is very small
 - Highly portable among many targets
 - But very limited in expressivity
- Accelerators can provide additional functionality
 - May not be portable between different targets
 - Under construction: library of standard accelerators

